

HOUSE PRICE PREDICTION PROJECT:

In []:

In [102]:

```
#Import required libraries:  
import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
import seaborn as sns  
  
#Ignore all warnings:  
import warnings  
warnings.filterwarnings("ignore")
```

In [103]:

```
#Import the dataset:  
data = pd.read_csv("dataset.csv")
```

get a general information about the dataset:

In [104]:

```
print("No. of rows X columns = ",data.shape)  
data.info()  
round(data.describe(),2)
```

```
No. of rows X columns = (14619, 23)  
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 14619 entries, 0 to 14618  
Data columns (total 23 columns):  
 #   Column                Non-Null Count  Dtype     
 ---  --        
 0   id                   14619 non-null   int64    
 1   Date                  14619 non-null   int64    
 2   number of bedrooms    14619 non-null   int64    
 3   number of bathrooms   14619 non-null   float64  
 4   living area           14619 non-null   int64    
 5   lot area              14619 non-null   int64    
 6   number of floors     14619 non-null   float64  
 7   waterfront present   14619 non-null   int64    
 8   number of views       14619 non-null   int64    
 9   condition of the house 14619 non-null   int64    
 10  grade of the house   14619 non-null   int64    
 11  Area of the house(excluding basement) 14619 non-null   int64    
 12  Area of the basement 14619 non-null   int64    
 13  Built Year            14619 non-null   int64    
 14  Renovation Year      14619 non-null   int64    
 15  Postal Code           14619 non-null   int64    
 16  Latitude               14619 non-null   float64  
 17  Longitude              14619 non-null   float64  
 18  living_area_renov     14619 non-null   int64    
 19  lot_area_renov        14619 non-null   int64    
 20  Number of schools nearby 14619 non-null   int64    
 21  Distance from the airport 14619 non-null   int64    
 22  Price                  14619 non-null   int64    
dtypes: float64(4), int64(19)  
memory usage: 2.6 MB
```

Out[104]:

	id	Date	number of bedrooms	number of bathrooms	living area	lot area	number of floors	waterfront present	number of views	condition of the house	... ::
count	1.461900e+04	14619.00	14619.00	14619.00	14619.00	14619.00	14619.00	14619.00	14619.00	14619.00	... 14
mean	6.762821e+09	42604.55	3.38	2.13	2098.16	15093.69	1.50	0.01	0.23	3.43	... 1
std	6.237160e+03	67.34	0.94	0.77	928.22	37920.89	0.54	0.09	0.77	0.66	...
min	6.762810e+09	42491.00	1.00	0.50	370.00	520.00	1.00	0.00	0.00	1.00	... 1
25%	6.762815e+09	42546.00	3.00	1.75	1440.00	5010.50	1.00	0.00	0.00	3.00	... 1
50%	6.762821e+09	42600.00	3.00	2.25	1930.00	7620.00	1.50	0.00	0.00	3.00	... 1
75%	6.762826e+09	42662.00	4.00	2.50	2570.00	10800.00	2.00	0.00	0.00	4.00	... 1
max	6.762832e+09	42734.00	33.00	8.00	13540.00	1074218.00	3.50	1.00	4.00	5.00	... 2

8 rows × 23 columns

In [105]:

```
print("Sample 5 entries of data: ")
data.sample(5)
```

Sample 5 entries of data:

Out[105]:

	id	Date	number of bedrooms	number of bathrooms	living area	lot area	number of floors	waterfront present	number of views	condition of the house	... ::	Built Year	Renovation Year
3338	6762826256	42543	4	1.50	1890	43560	1.0	0	0	4	... 1974	0	
7456	6762818993	42602	3	2.00	1500	2500	2.0	0	0	3	... 2002	0	
10164	6762830891	42648	3	1.00	1090	17630	1.0	0	0	4	... 1962	0	
12782	6762813696	42693	4	2.50	2370	2971	2.0	0	2	3	... 2008	0	
10668	6762813432	42657	4	2.75	2660	4500	1.5	0	0	5	... 1909	0	

5 rows × 23 columns

In [106]:

```
print("First 5 entries of data: ")
data.head()
```

First 5 entries of data:

Out[106]:

	id	Date	number of bedrooms	number of bathrooms	living area	lot area	number of floors	waterfront present	number of views	condition of the house	... ::	Built Year	Renovation Year	Po C
0	6762810635	42491	4	2.50	2920	4000	1.5	0	0	5	... 1909	0	122	
1	6762810998	42491	5	2.75	2910	9480	1.5	0	0	3	... 1939	0	122	
2	6762812605	42491	4	2.50	3310	42998	2.0	0	0	3	... 2001	0	122	
3	6762812919	42491	3	2.00	2710	4500	1.5	0	0	4	... 1929	0	122	
4	6762813105	42491	3	2.50	2600	4750	1.0	0	0	4	... 1951	0	122	

5 rows × 23 columns

In [107]:

```
print("Last 5 entries of data: ")
data.tail()
```

Last 5 entries of data:

Out[107]:

	id	Date	number of bedrooms	number of bathrooms	living area	lot area	number of floors	waterfront present	number of views	condition of the house	...	Built Year	Renovation Year
14614	6762830250	42734	2	1.5	1556	20000	1.0	0	0	4	...	1957	0
14615	6762830339	42734	3	2.0	1680	7000	1.5	0	0	4	...	1968	0
14616	6762830618	42734	2	1.0	1070	6120	1.0	0	0	3	...	1962	0
14617	6762830709	42734	4	1.0	1030	6621	1.0	0	0	4	...	1955	0
14618	6762831463	42734	3	1.0	900	4770	1.0	0	0	3	...	1969	2009

5 rows x 23 columns

In [108]:

```
#Check for any null-values in each column:
result = data.isna().sum().reset_index()
result.columns=["Column-name","No. of null values"]
print(result)

print("\n\n\n")

#check for any duplicated values:
result = data.duplicated().sum()
print("No. of duplicated values = ",result)
```

	Column-name	No. of null values
0	id	0
1	Date	0
2	number of bedrooms	0
3	number of bathrooms	0
4	living area	0
5	lot area	0
6	number of floors	0
7	waterfront present	0
8	number of views	0
9	condition of the house	0
10	grade of the house	0
11	Area of the house(excluding basement)	0
12	Area of the basement	0
13	Built Year	0
14	Renovation Year	0
15	Postal Code	0
16	Latitude	0
17	Longitude	0
18	living_area_renov	0
19	lot_area_renov	0
20	Number of schools nearby	0
21	Distance from the airport	0
22	Price	0

No. of duplicated values = 0

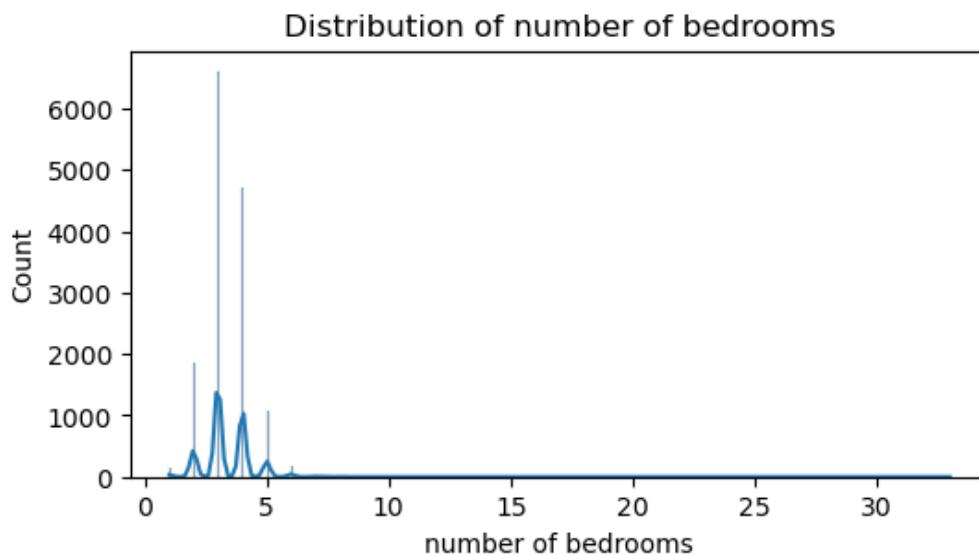
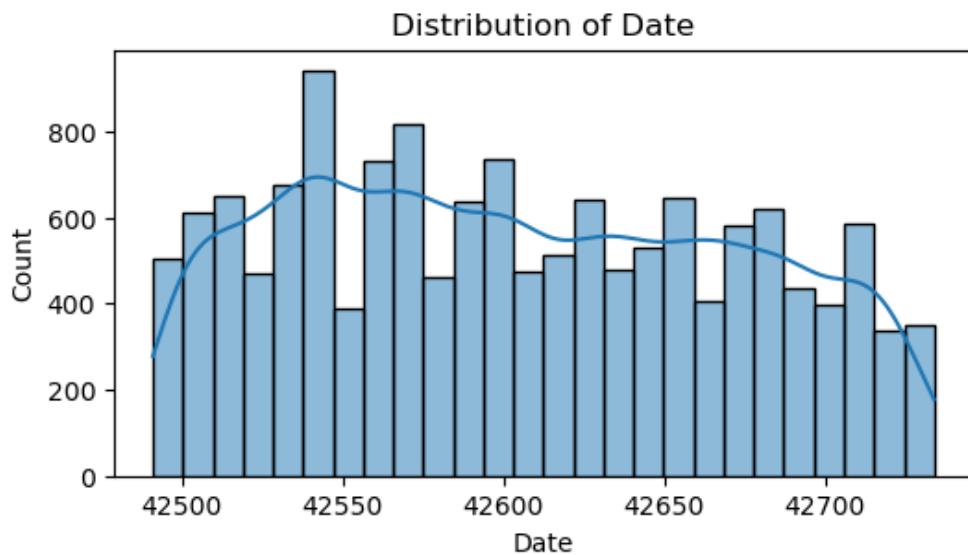
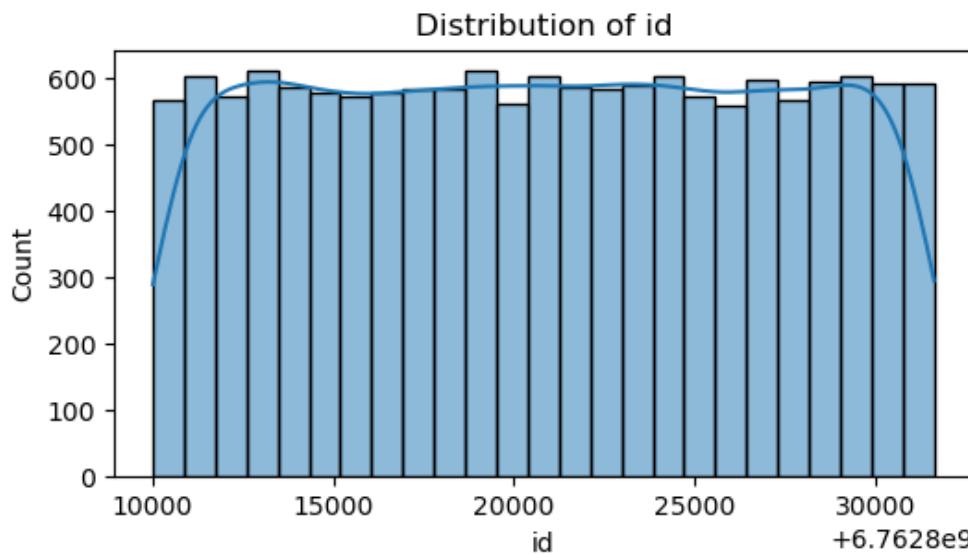
In []:

EDA:

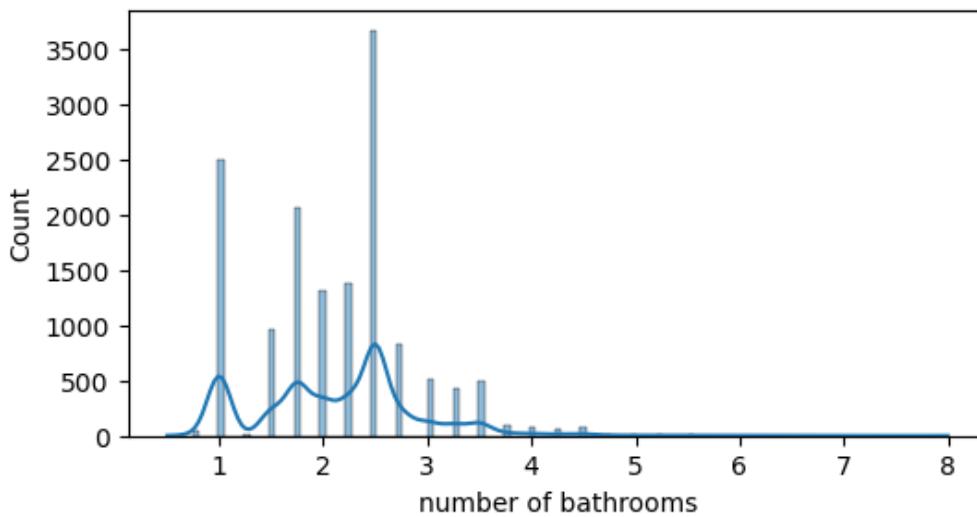
In [109]:

```
numerical_cols = data.select_dtypes(include=['int64', 'float64']).columns

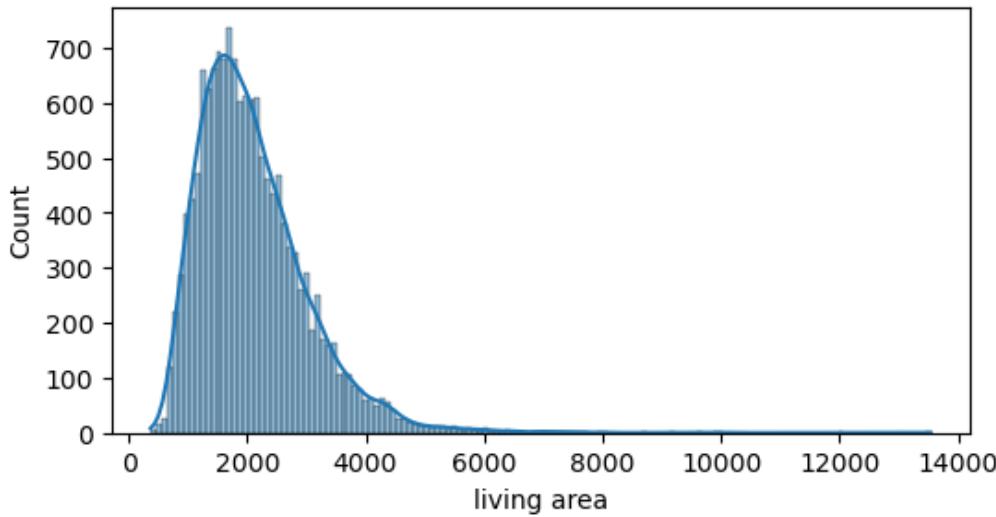
for col in numerical_cols:
    plt.figure(figsize=(6,3))
    sns.histplot(data[col], kde=True)
    plt.title(f"Distribution of {col}")
    plt.show()
```



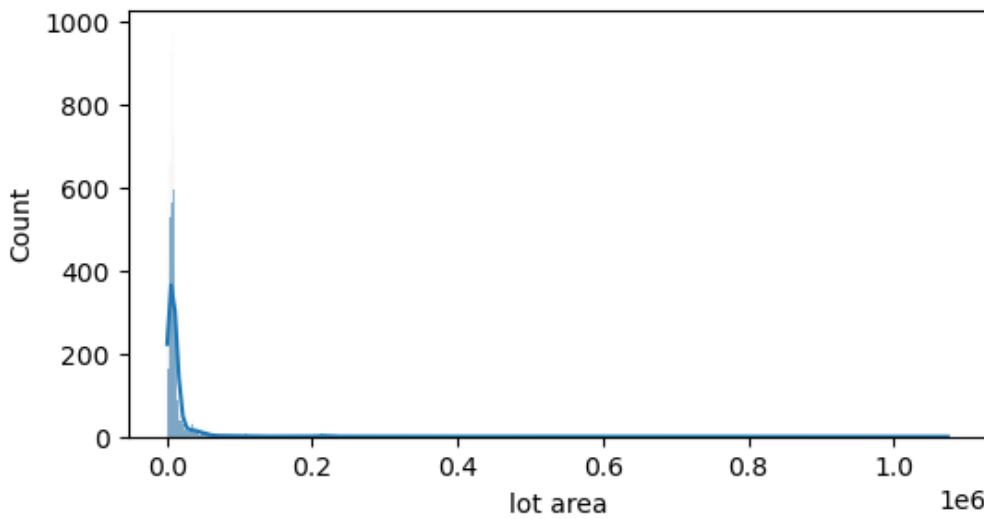
Distribution of number of bathrooms



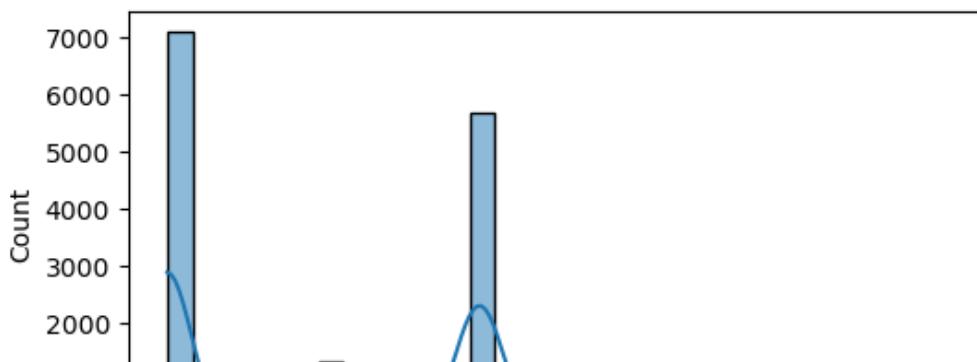
Distribution of living area

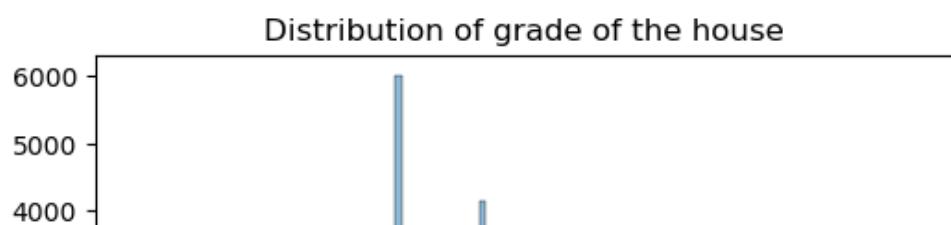
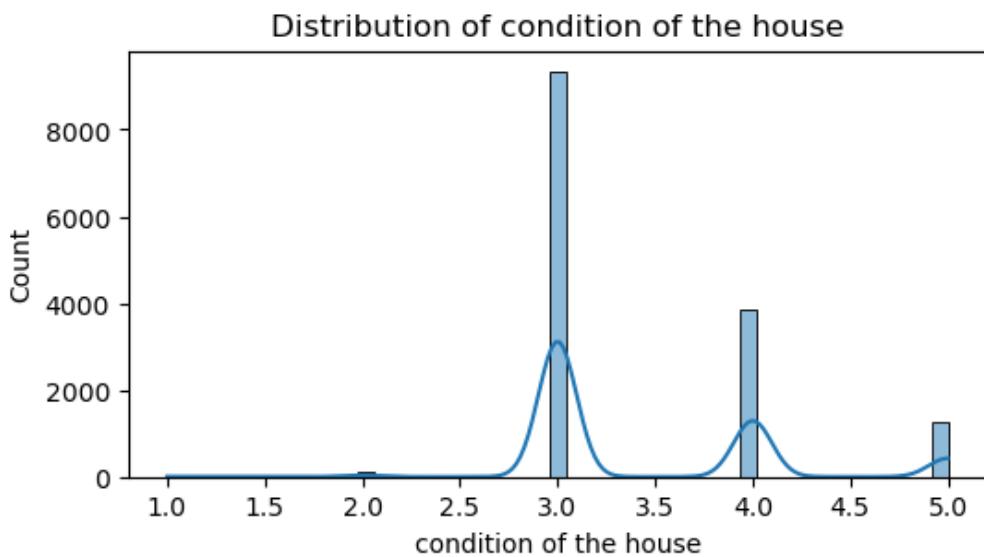
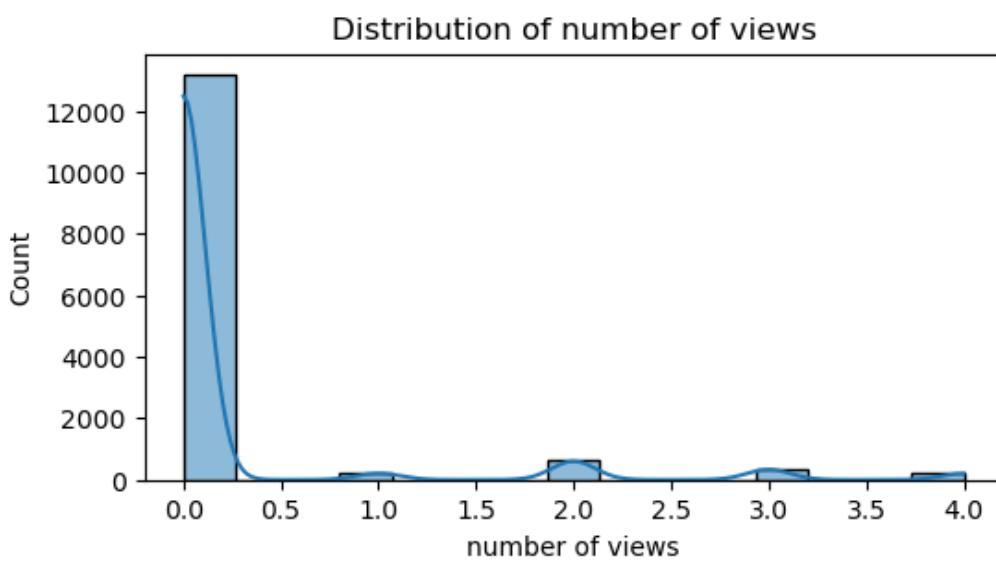
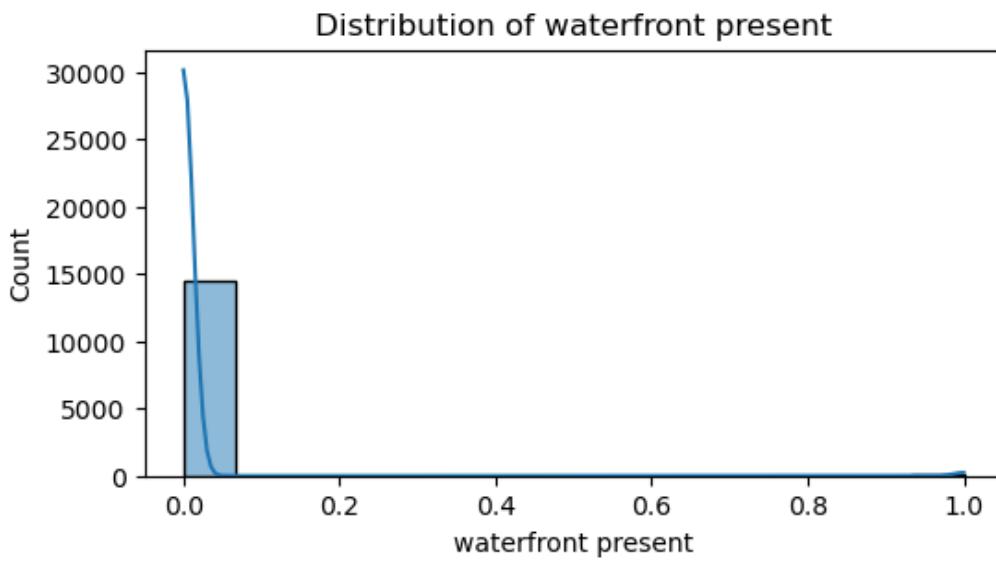
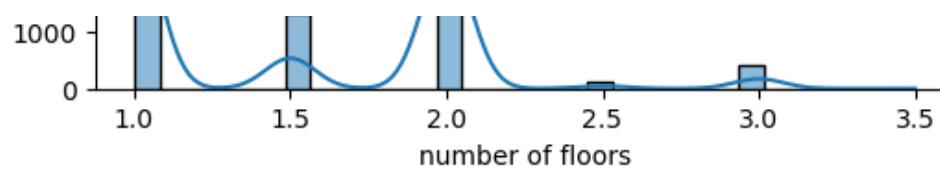


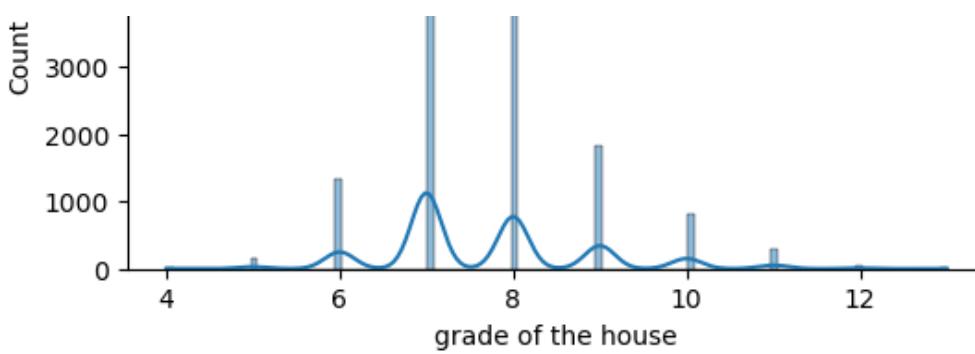
Distribution of lot area



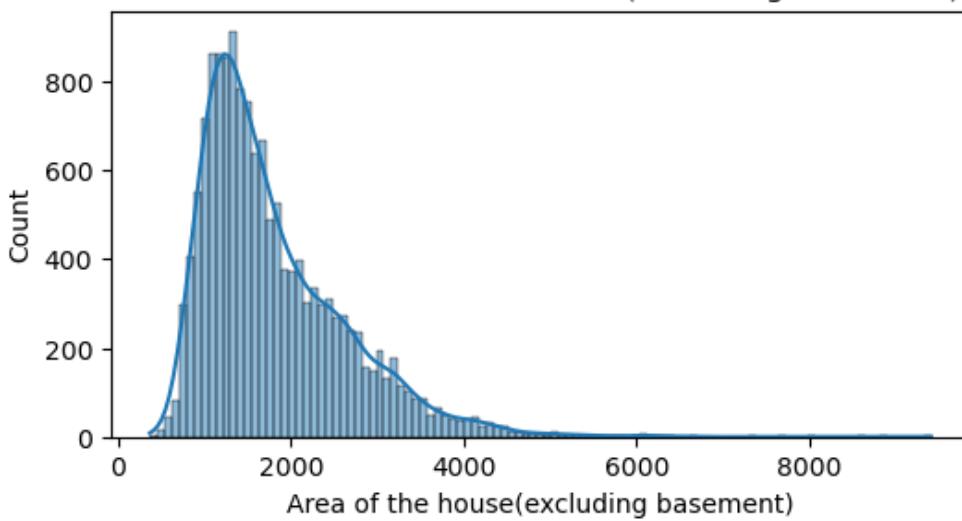
Distribution of number of floors



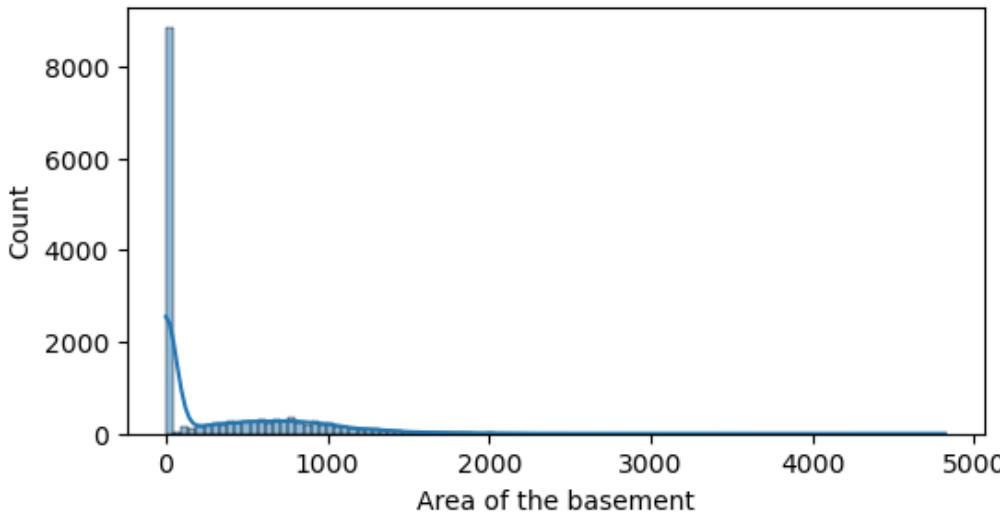




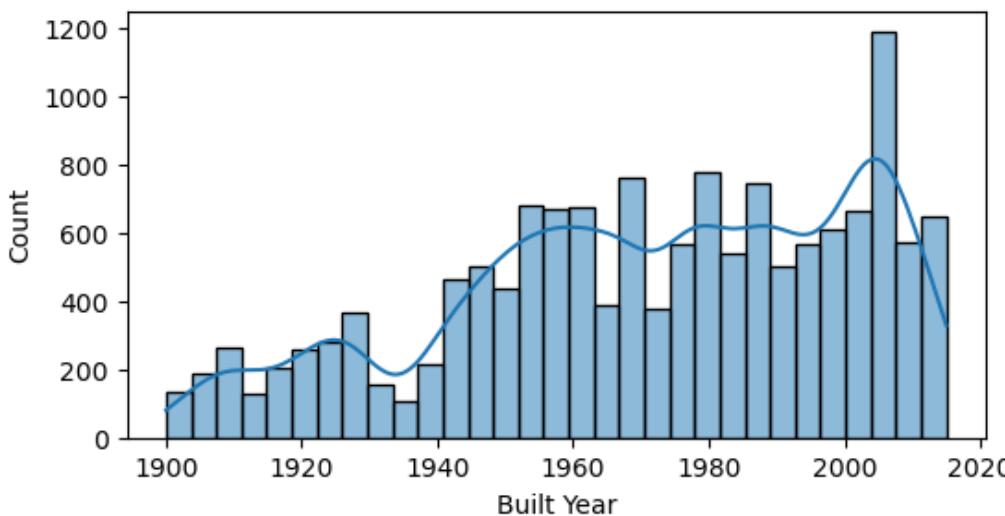
Distribution of Area of the house(excluding basement)



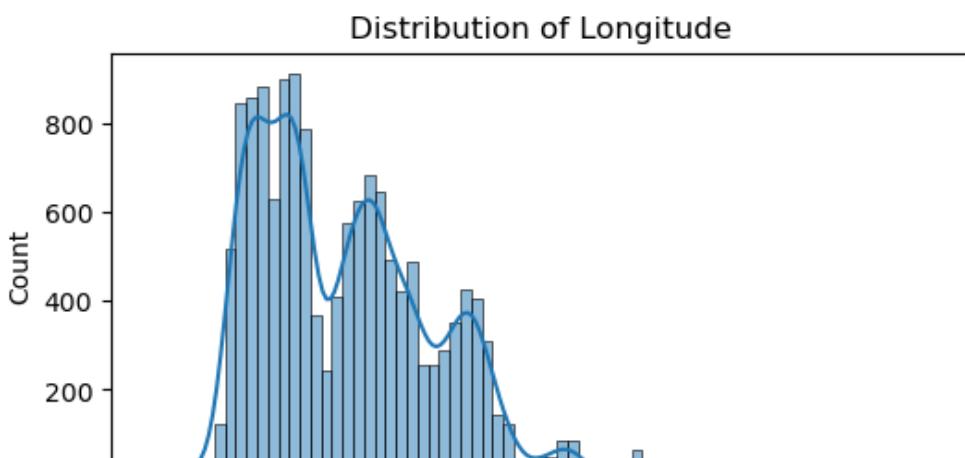
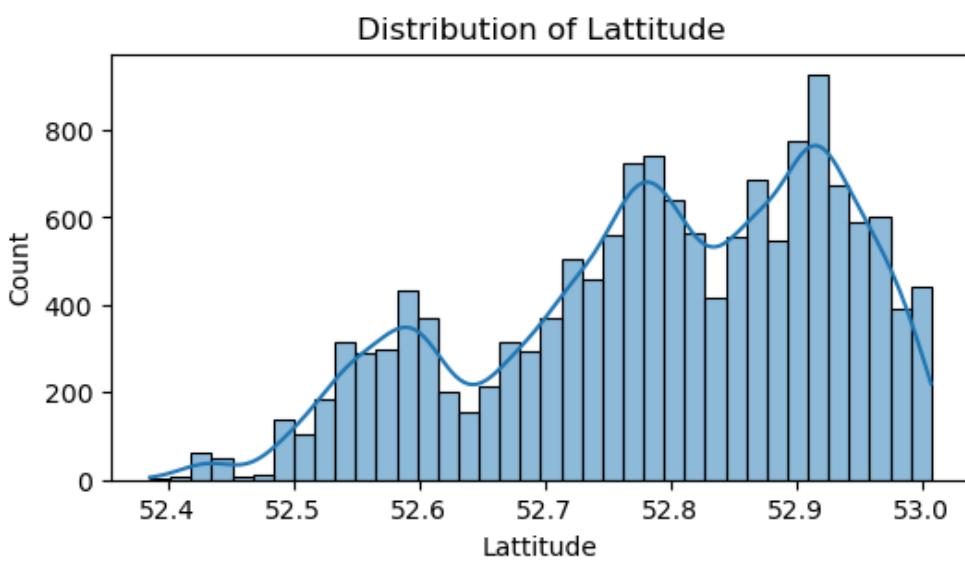
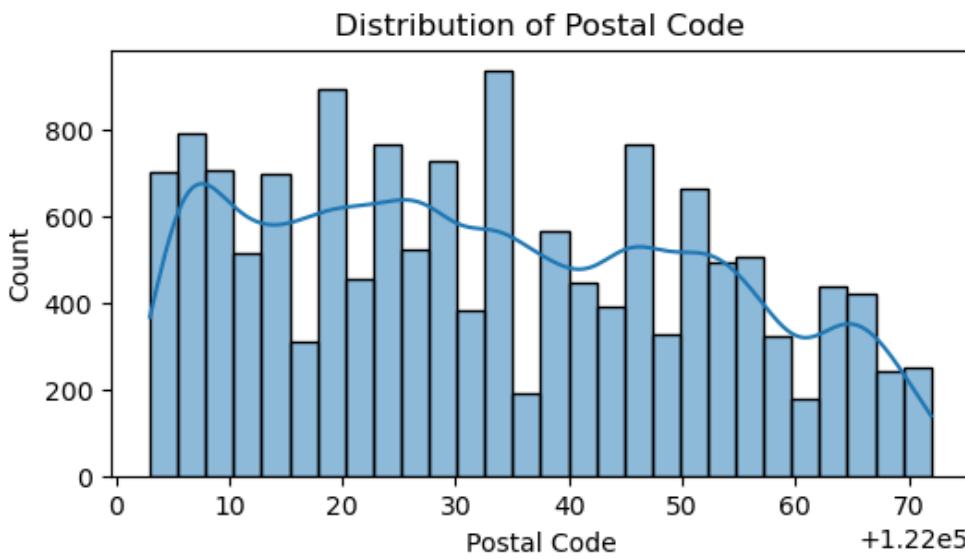
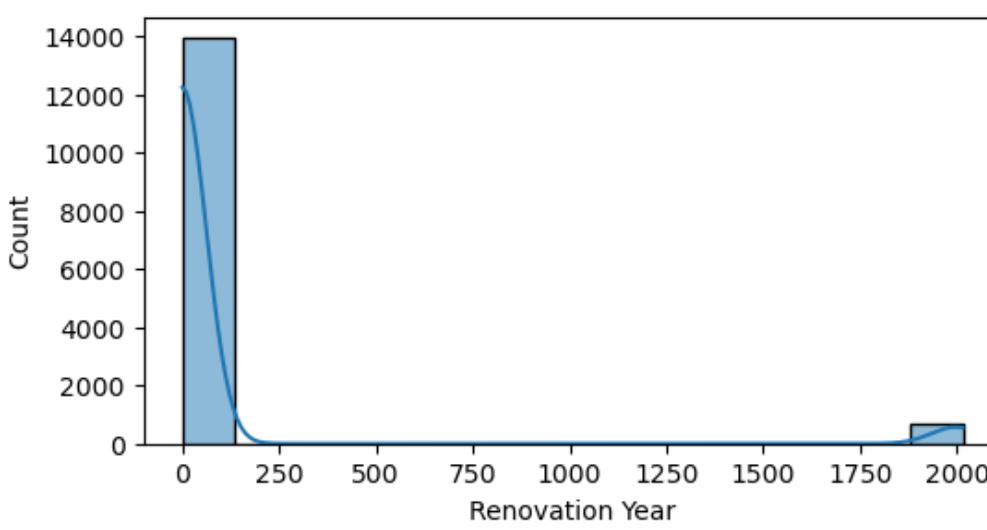
Distribution of Area of the basement

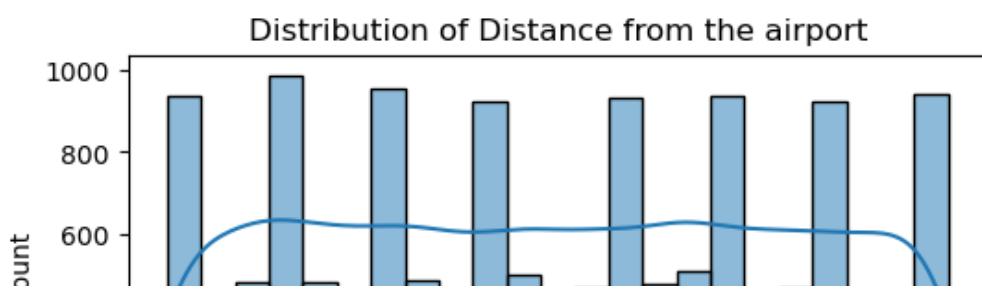
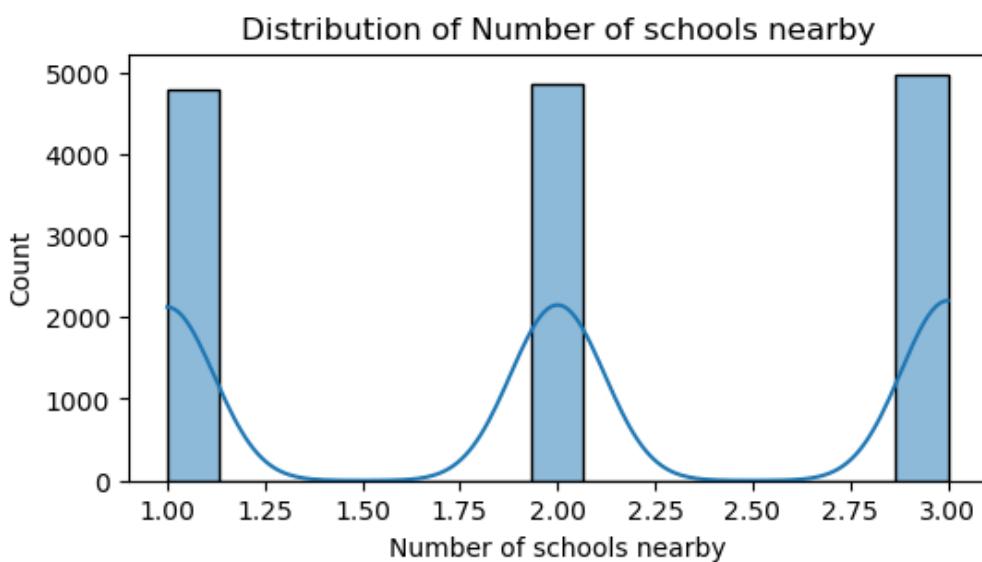
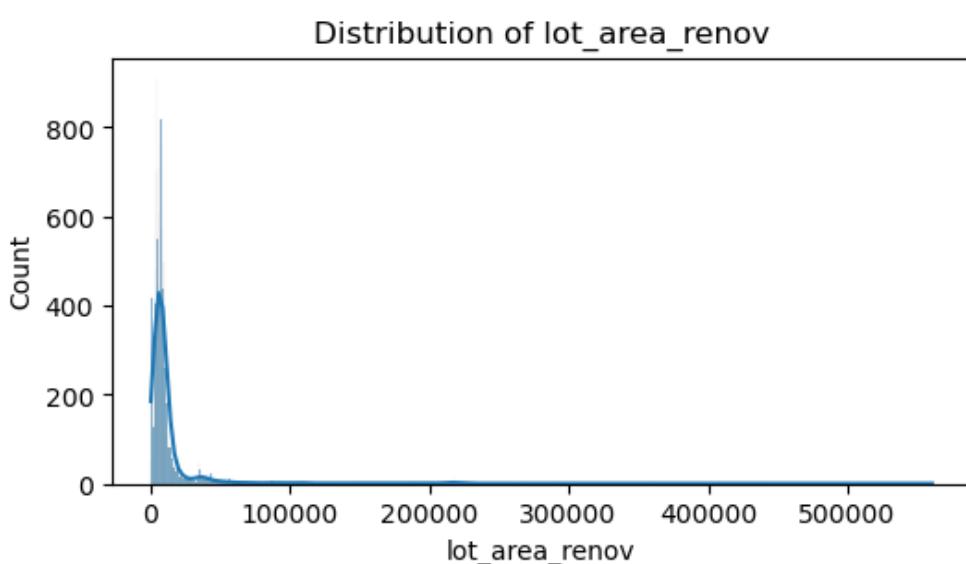
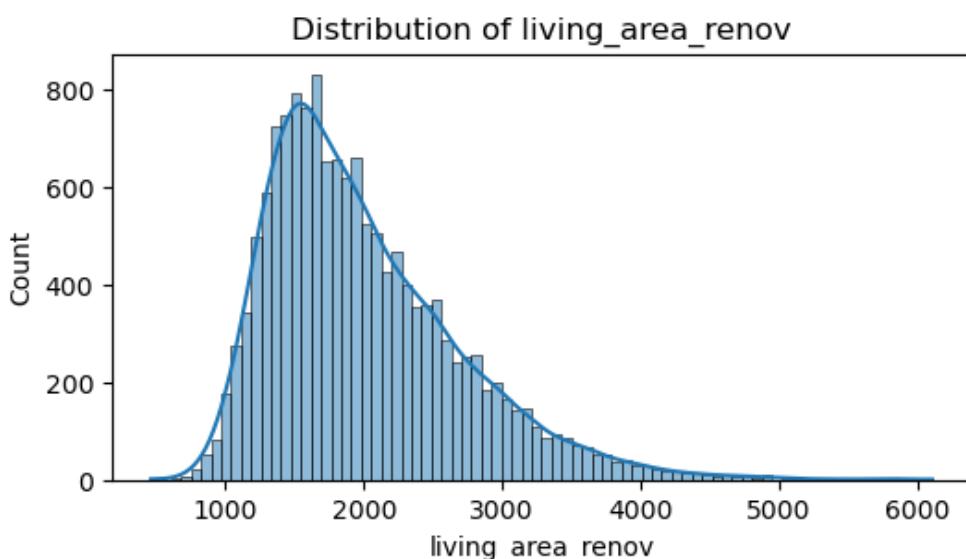
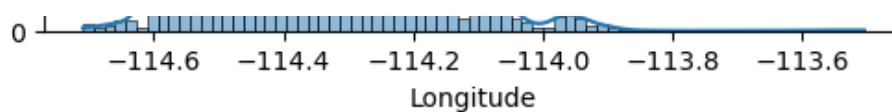


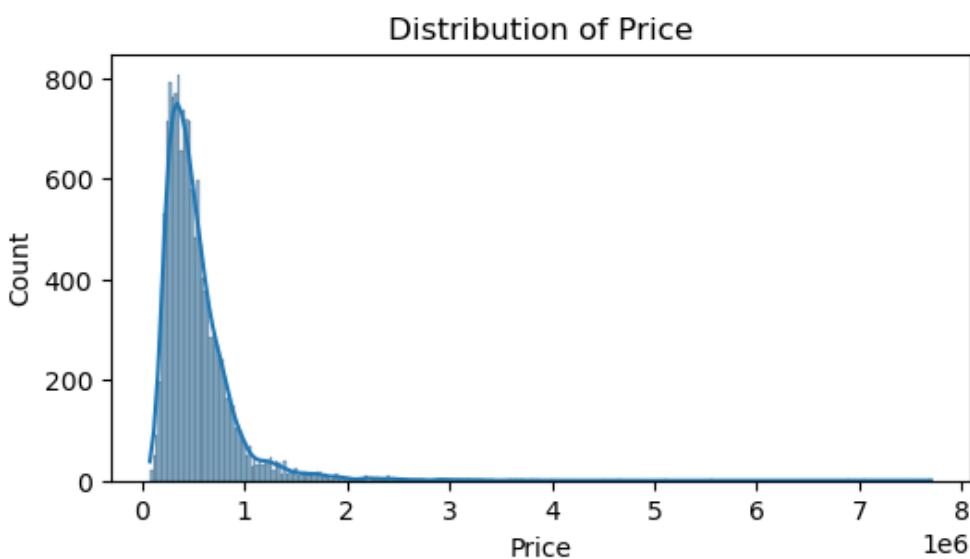
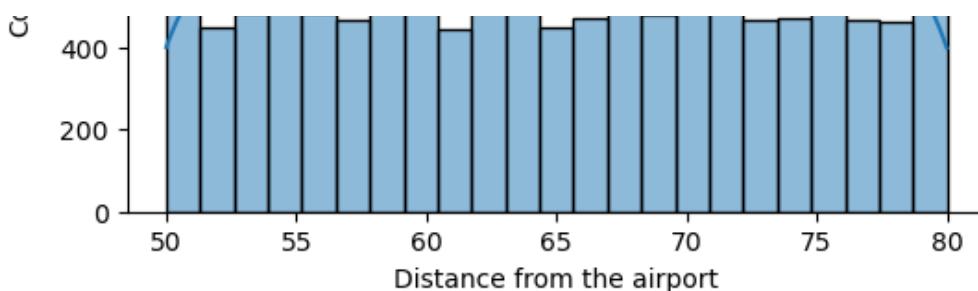
Distribution of Built Year



Distribution of Renovation Year

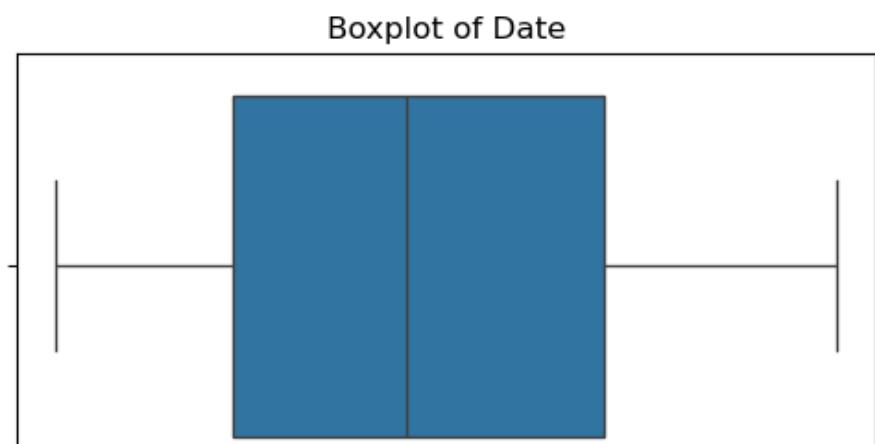
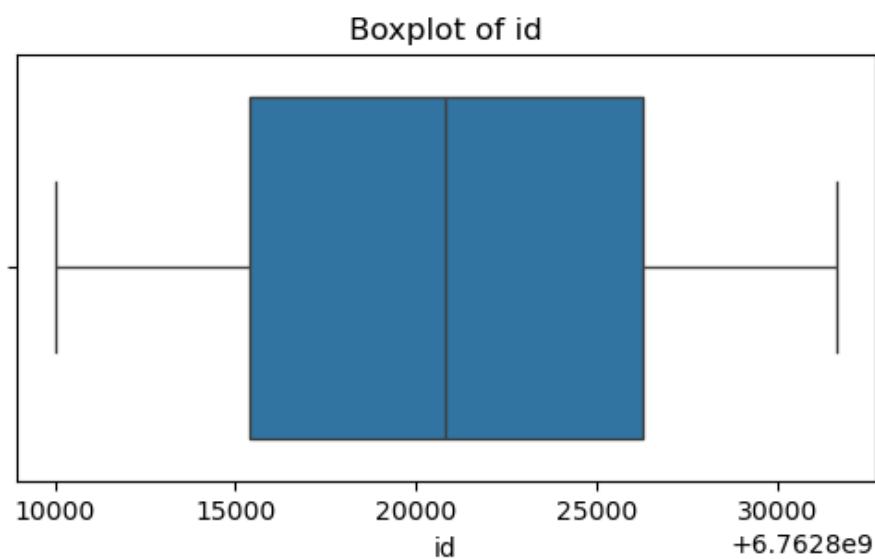


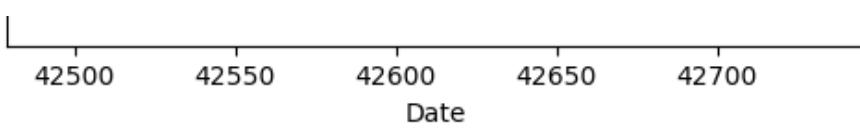




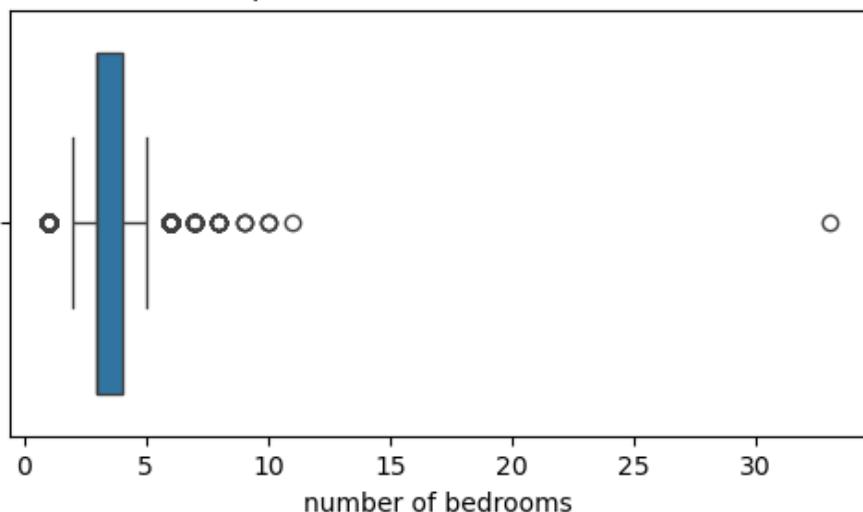
In [110]:

```
for col in numerical_cols:  
    plt.figure(figsize=(6,3))  
    sns.boxplot(x=data[col])  
    plt.title(f"Boxplot of {col}")  
    plt.show()
```

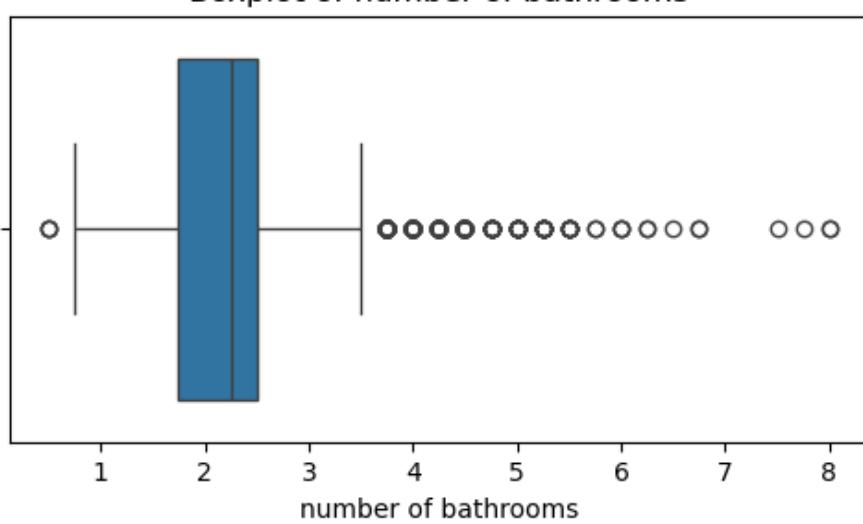




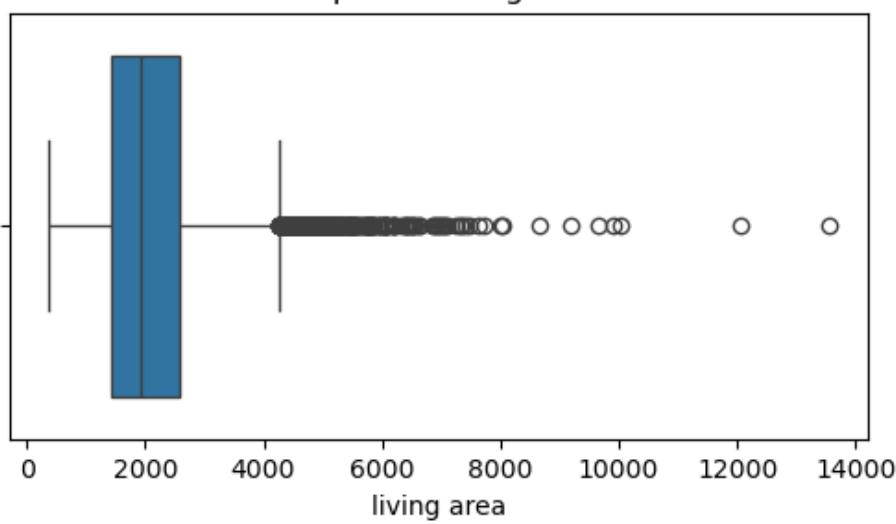
Boxplot of number of bedrooms



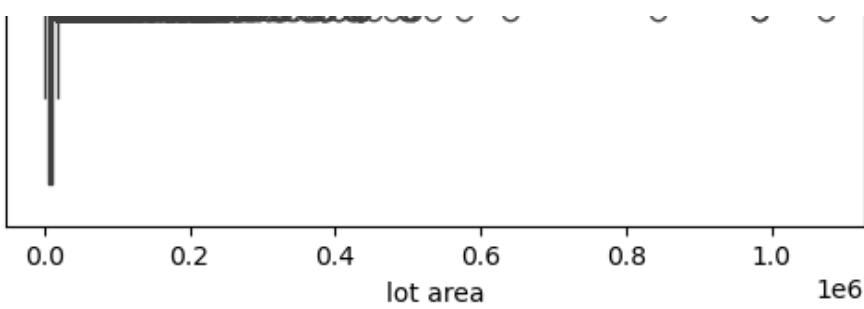
Boxplot of number of bathrooms



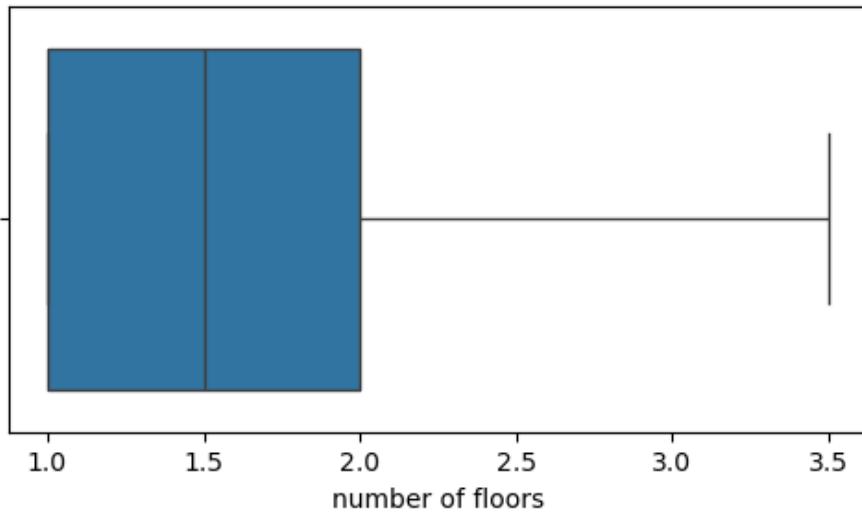
Boxplot of living area



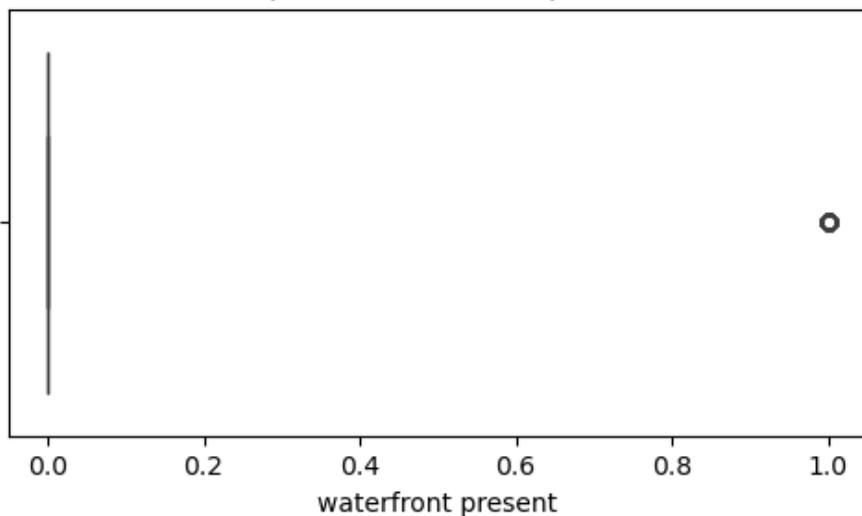
Boxplot of lot area



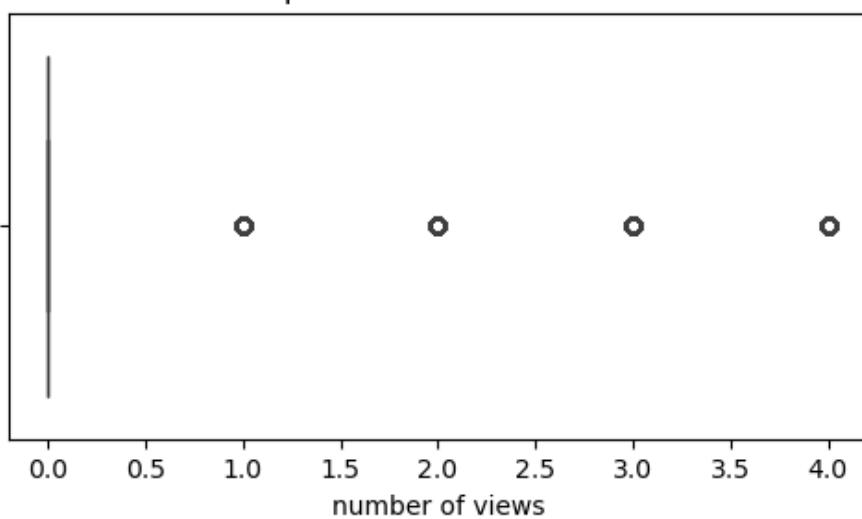
Boxplot of number of floors



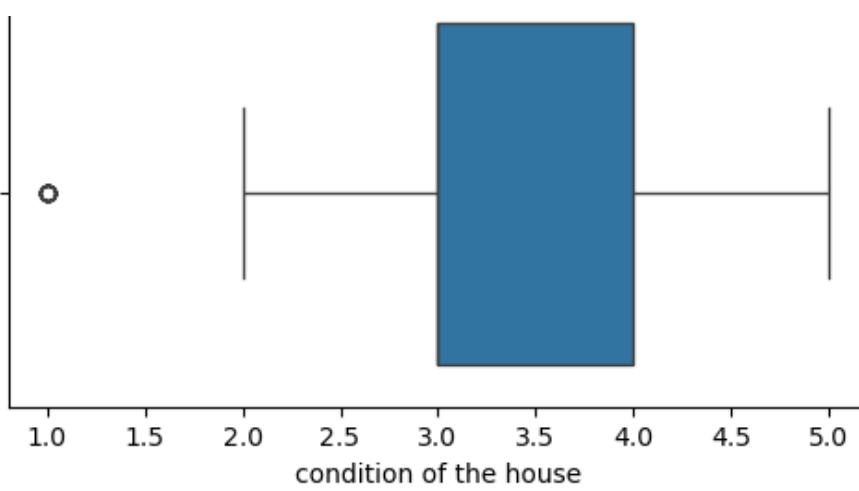
Boxplot of waterfront present



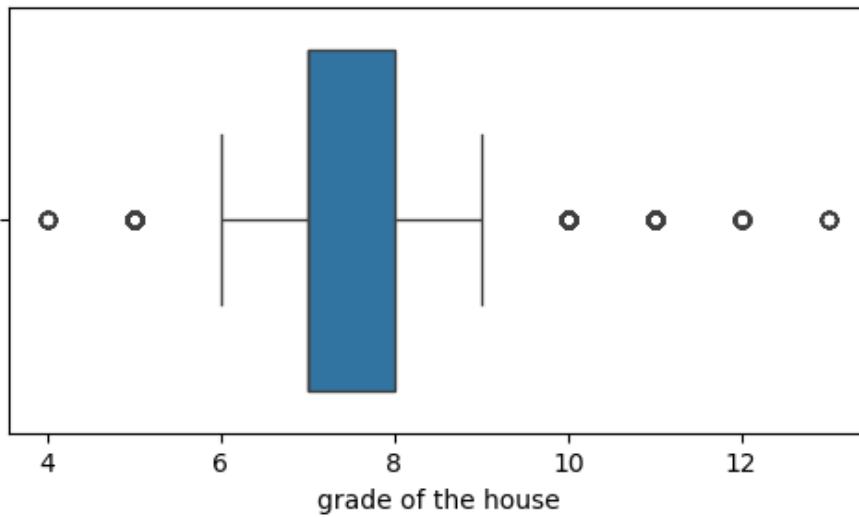
Boxplot of number of views



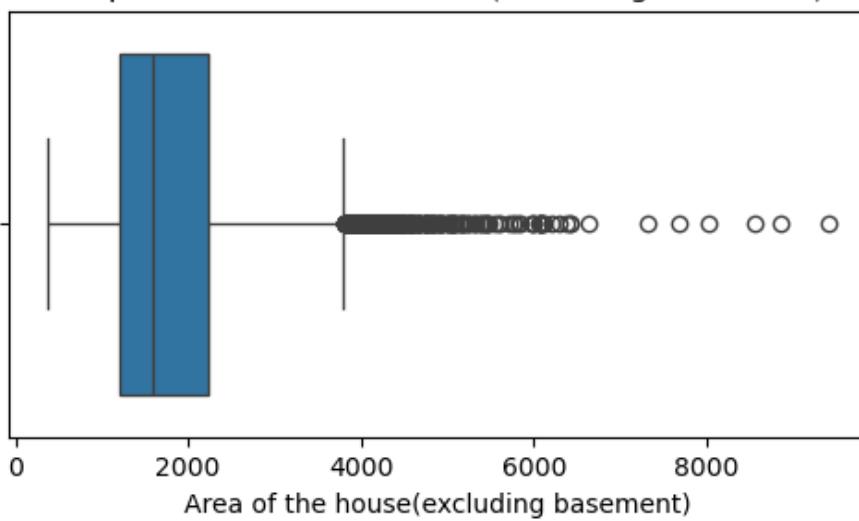
Boxplot of condition of the house



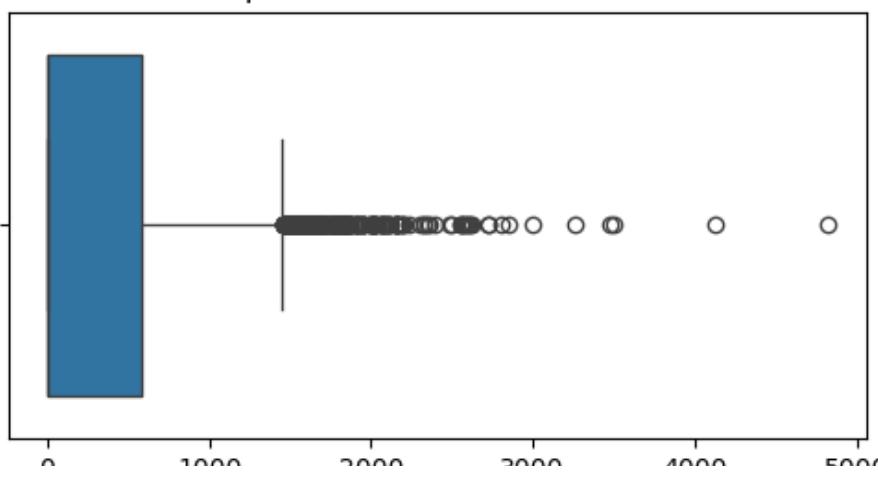
Boxplot of grade of the house



Boxplot of Area of the house(excluding basement)

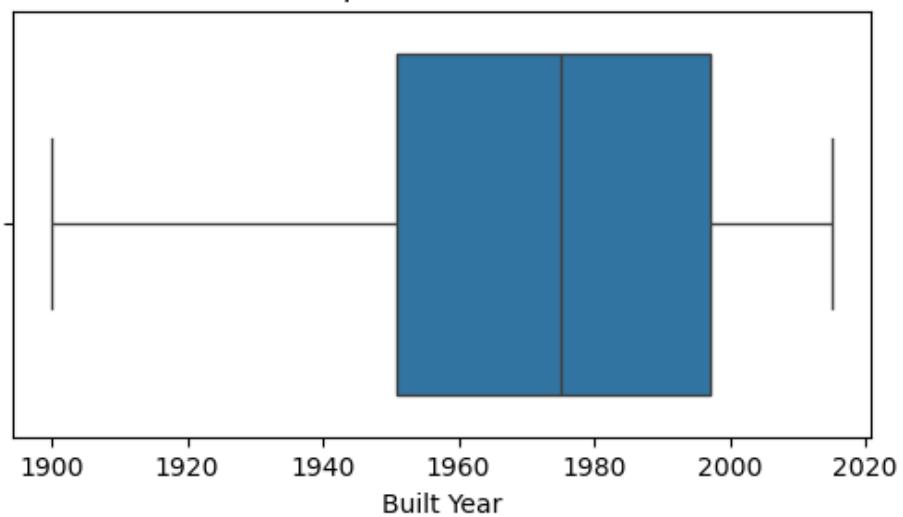


Boxplot of Area of the basement

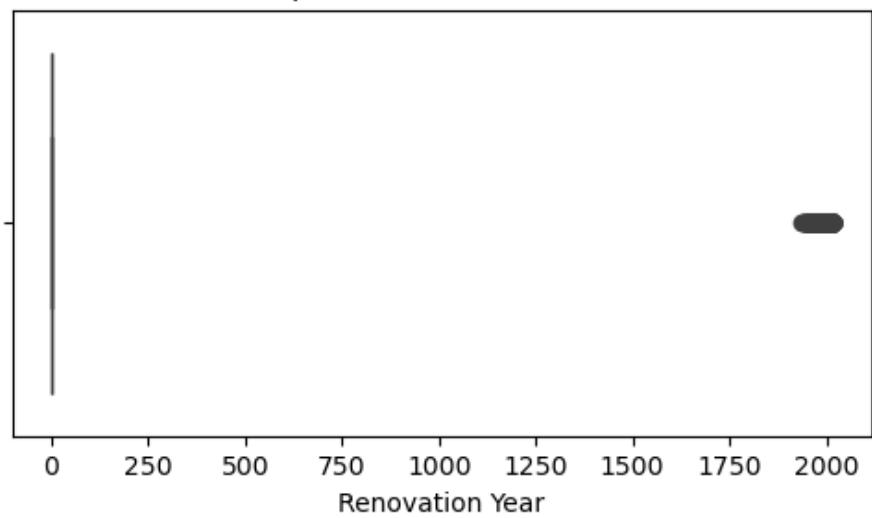


0 1000 2000 3000 4000 5000
Area of the basement

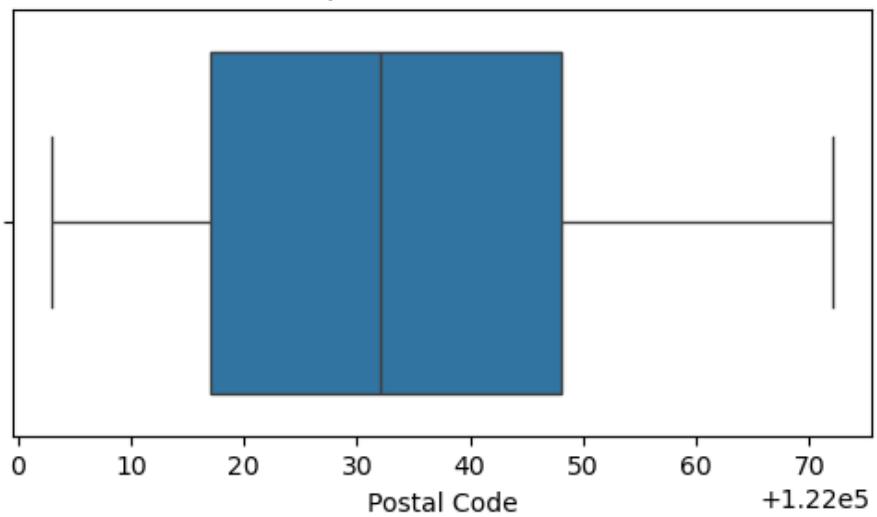
Boxplot of Built Year



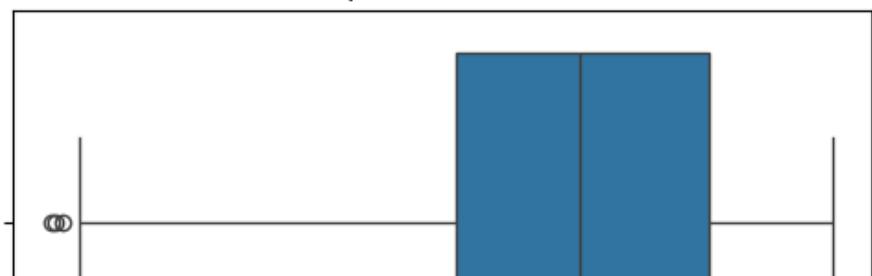
Boxplot of Renovation Year

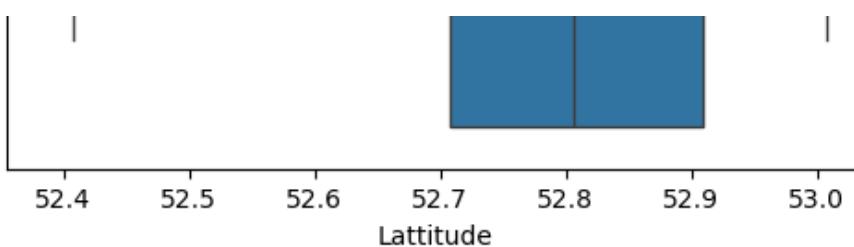


Boxplot of Postal Code

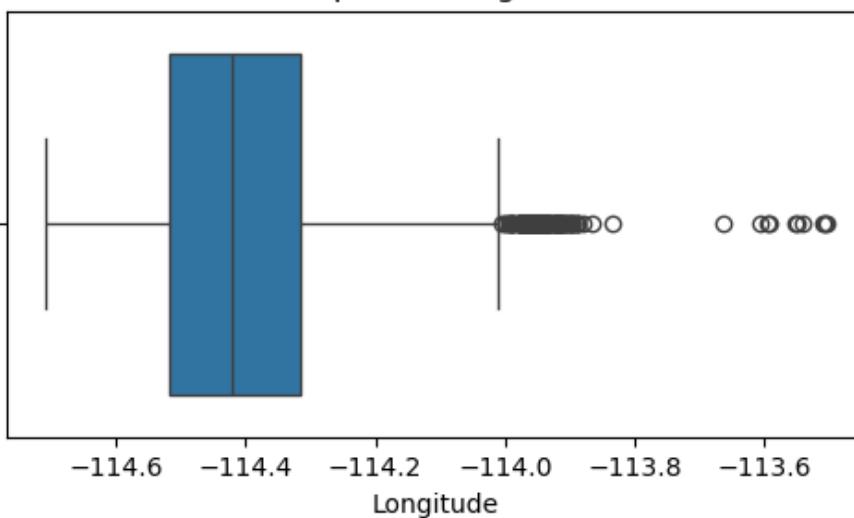


Boxplot of Latitude

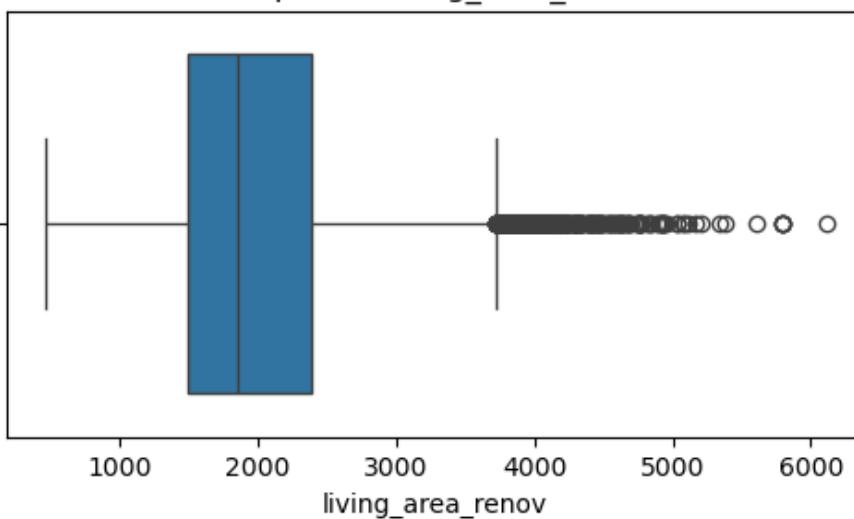




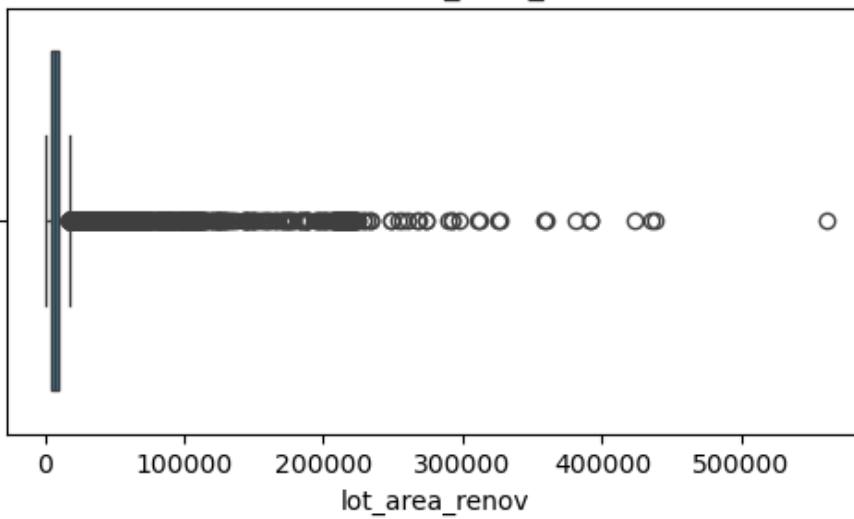
Boxplot of Longitude



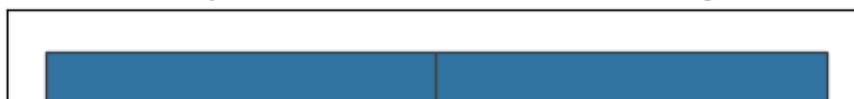
Boxplot of living_area_renov

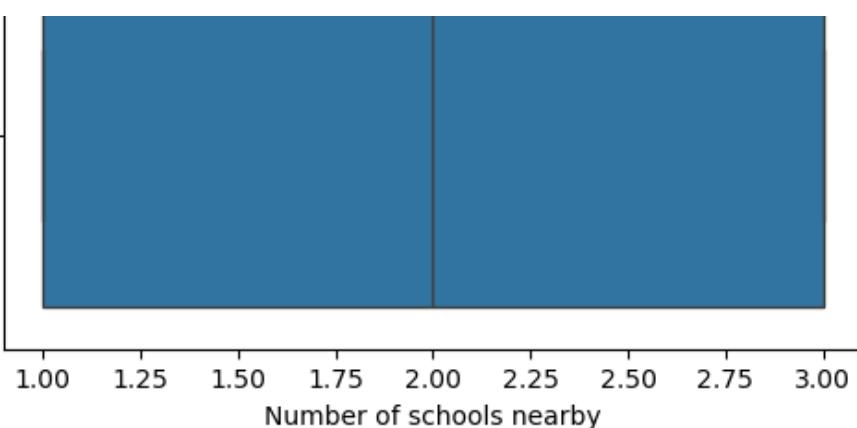


Boxplot of lot_area_renov

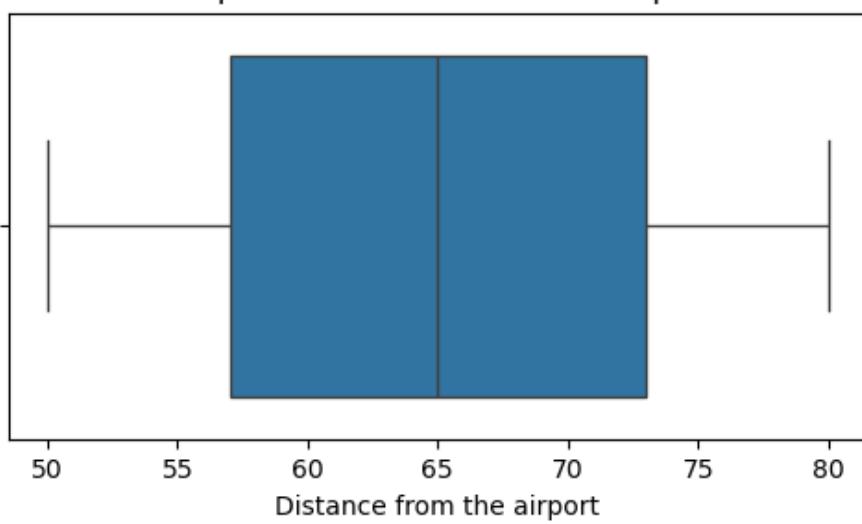


Boxplot of Number of schools nearby

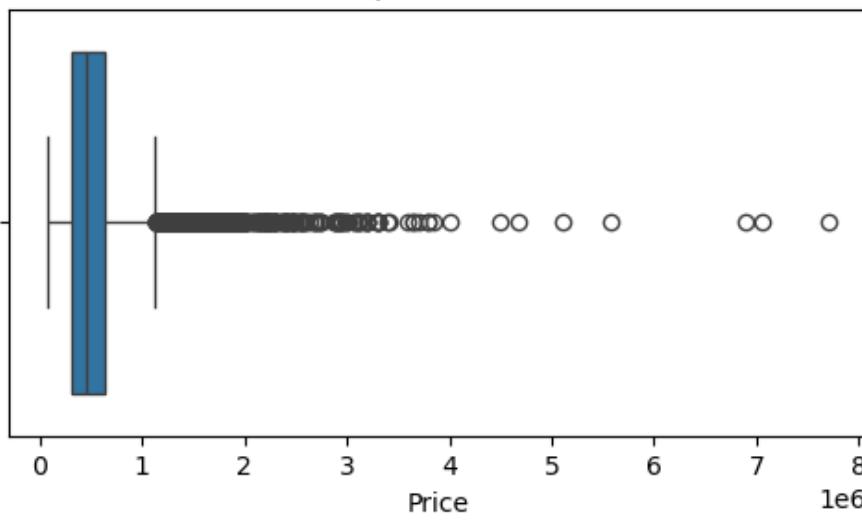




Boxplot of Distance from the airport



Boxplot of Price



Now check the relationship between:

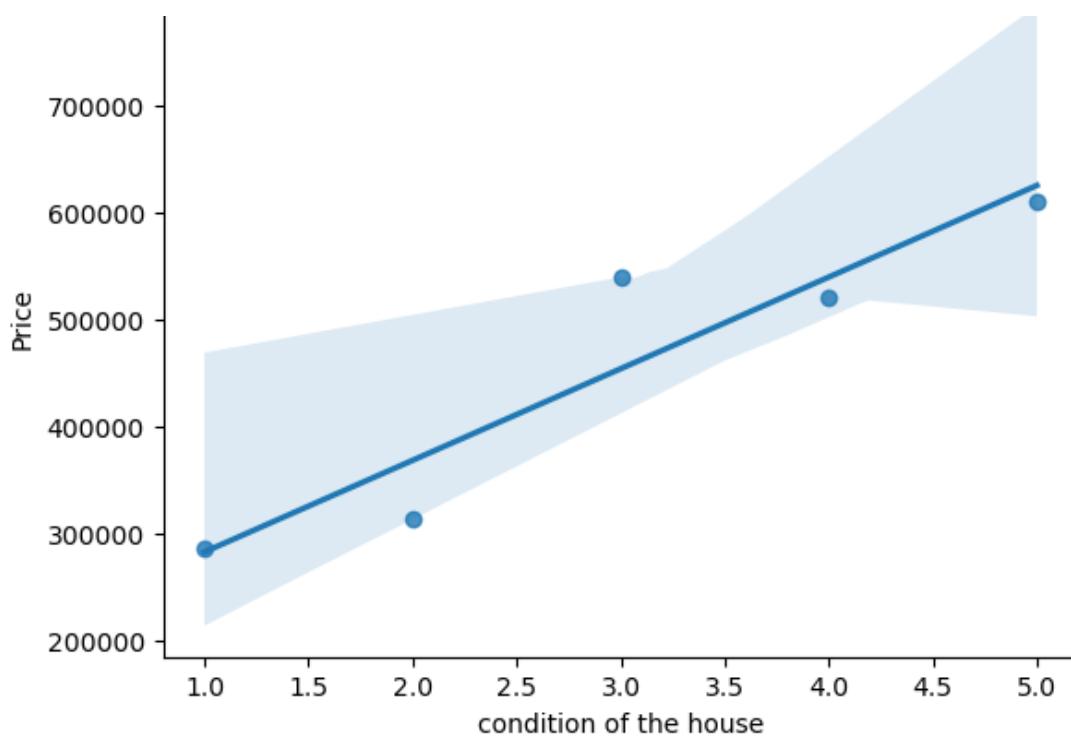
- Condition of the house
- Price of the house (on average)

In [111]:

```
df = data.groupby("condition of the house") ["Price"].mean().reset_index().sort_values(by="Price")
sns.regplot(x="condition of the house",y="Price",data=df)
```

Out[111]:

```
<Axes: xlabel='condition of the house', ylabel='Price'>
```



Conclusion at this point:

- Houses in proper condition are costly.
 - Whereas; houses (not in proper condition) are cheap.

In [112]:

```
data.head()
```

Out [112] :

	id	Date	number of bedrooms	number of bathrooms	living area	lot area	number of floors	waterfront present	number of views	condition of the house	...	Built Year	Renovation Year	PoC
0	6762810635	42491	4	2.50	2920	4000	1.5	0	0	5	...	1909	0	122
1	6762810998	42491	5	2.75	2910	9480	1.5	0	0	3	...	1939	0	122
2	6762812605	42491	4	2.50	3310	42998	2.0	0	0	3	...	2001	0	122
3	6762812919	42491	3	2.00	2710	4500	1.5	0	0	4	...	1929	0	122
4	6762813105	42491	3	2.50	2600	4750	1.0	0	0	4	...	1951	0	122

5 rows × 23 columns

ML QUESTIONS:

Q1] Predict house-prices based on:

- Number of bedrooms
 - Numbers of bathrooms
 - Living area
 - Condition of the house

In [113]:

```
x = data[["number of bedrooms", "number of bathrooms", "living area", "condition of the house"]]
```

```
y = data[["Price"]]
```

In [114]:

```
from sklearn.model_selection import train_test_split, GridSearchCV
X_train , X_test , y_train, y_test = train_test_split(X,y,test_size=0.2)

from sklearn.tree import DecisionTreeRegressor
param_grid = {
    "criterion" : ["squared_error", "friedman_mse", "absolute_error"],
    "splitter" : ["best", "random"],
    "max_depth": [None, 10,20,30,40,50],
    "min_samples_split": [2,5,10],
    "min_samples_leaf": [1,2,4]
}

tree_model = DecisionTreeRegressor()
grid_tree= GridSearchCV(estimator = tree_model, param_grid = param_grid)
grid_tree.fit(X_train, y_train)
```

Out[114]:

```
▶      GridSearchCV
          i ?
▶      best_estimator_:
DecisionTreeRegressor
▶      DecisionTreeRegressor
          ?
```

In [115]:

```
#Check the performance of this model:

##but first make the predictions for the testing-dataset:
tree_preds = grid_tree.predict(X_test)
print(tree_preds)

from sklearn.metrics import mean_squared_error
mean_squared_error(y_test, tree_preds)
```

```
[355495.67315175 522026.90358744 238611.75555556 ... 353295.83333333
 539271.16622691 421135.27925532]
```

Out[115]:

```
61623684399.663
```

In [116]:

```
##Now make the predictions and check the performance of the linear-regression model:
from sklearn.linear_model import LinearRegression
lr = LinearRegression()
lr.fit(X_train, y_train)

predslr = lr.predict(X_test) #predictions for the testing data using linear-regression
from sklearn.metrics import mean_squared_error
mean_squared_error(y_test, predslr)
```

Out[116]:

```
60576794781.62602
```

In [117]:

```
## Now for both the models check the mean_absolute_error values:
from sklearn.metrics import mean_absolute_error
print("For GridSearchCV: ",mean_absolute_error(y_test, tree_preds))
print("For LinearRegression: ",mean_absolute_error(y_test, predslr))
```

```
For GridSearchCV: 161548.44850459037  
For LinearRegression: 165031.71973815977
```

In [118]:

```
#Now make the predictions using Random-forest model:  
from sklearn.ensemble import RandomForestRegressor  
rfrmodel = RandomForestRegressor()  
param_gridrfr = {  
    "max_depth": [5,10,15],  
    "n_estimators": [2,3,4,5,6,7,8,9,10]  
}  
  
gridrfr = GridSearchCV(rfrmodel,param_gridrfr)  
gridrfr.fit(X_train, y_train)
```

Out[118]:

```
►      GridSearchCV  
          i ?  
►      best_estimator_:  
          RandomForestRegressor  
►      RandomForestRegressor  
          ?
```

In [119]:

```
rfrpredictions = gridrfr.predict(X_test)  
mean_absolute_error(y_test,rfrpredictions)
```

Out[119]:

```
157897.00309501652
```

In [121]:

```
### Now for all the models check the mean_absolute_error values:  
from sklearn.metrics import mean_absolute_error  
print("For GridSearchCV: ",mean_absolute_error(y_test, tree_preds))  
print("For LinearRegression: ",mean_absolute_error(y_test, predslr))  
print("For Random-Forest: ",mean_absolute_error(y_test, rfrpredictions))
```

```
For GridSearchCV: 161548.44850459037  
For LinearRegression: 165031.71973815977  
For Random-Forest: 157897.00309501652
```

In [122]:

```
#Save your random-forest model on your system:  
import joblib  
joblib.dump(gridrfr,"model.pkl")
```

Out[122]:

```
['model.pkl']
```

In []: