

LAPTOP PRICE PREDICTOR:

In []:

Ignore all warnings :

In [567]:

```
import warnings  
warnings.filterwarnings('ignore')
```

In []:

LOAD SOME LIBRARIES INITIALLY:

In [568]:

```
import numpy as np  
import pandas as pd  
  
import seaborn as sns  
import matplotlib.pyplot as plt
```

In []:

LOAD THE DATASET:

In [569]:

```
df = pd.read_csv("laptop_data.csv")
```

In [570]:

```
df.head()
```

Out[570]:

| | Unnamed: 0 | Company | TypeName | Inches | ScreenResolution | Cpu | Ram | Memory | Gpu | OpSys | Weight | Price |
|---|---------------|---------|-----------|--------|--|-------------------------------------|------|---------------------------|---------------------------------------|-------|--------|-----------|
| 0 | 0 | Apple | Ultrabook | 13.3 | IPS Panel Retina Display 2560x1600 | Intel Core i5 2.3GHz | 8GB | 128GB SSD | Intel Iris Plus Graphics 640 | macOS | 1.37kg | 71378.68 |
| 1 | 1 | Apple | Ultrabook | 13.3 | 1440x900 | Intel Core i5 1.8GHz | 8GB | 128GB Flash Storage | Intel HD Graphics 6000 | macOS | 1.34kg | 47895.52 |
| 2 | 2 | HP | Notebook | 15.6 | Full HD 1920x1080 | Intel Core i5 7200U 2.5GHz | 8GB | 256GB SSD | Intel HD Graphics 620 | No OS | 1.86kg | 30636.00 |
| 3 | 3 | Apple | Ultrabook | 15.4 | IPS Panel Retina Display 2880x1800 | Intel Core i7 2.7GHz | 16GB | 512GB SSD | AMD Radeon Pro 455 | macOS | 1.83kg | 135195.33 |

| Unnamed: 4 | Company Apple | Type Name Ultrabook | Inches 13.3 | Screen Resolution Display | Opal Core i5 | Ram 8GB | Memory 256GB | Intel Iris Gpu Plus SSD | OpSys macOS | Weight 1.37kg | Price 96095.80\$ |
|---------------|------------------|------------------------|----------------|------------------------------|-----------------|------------|-----------------|-------------------------------|----------------|------------------|---------------------|
| | | | | 2560x1600 | 3.1GHz | | | 650 | | | |

In []:

DETAILED INFORMATION ABOUT THE DATASET:

In [571]:

```
df.duplicated().sum()
```

Out[571]:

```
np.int64(0)
```

In [572]:

```
df.shape
```

Out[572]:

```
(1303, 12)
```

In [573]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1303 entries, 0 to 1302
Data columns (total 12 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Unnamed: 0        1303 non-null   int64  
 1   Company          1303 non-null   object  
 2   TypeName         1303 non-null   object  
 3   Inches           1303 non-null   float64
 4   ScreenResolution 1303 non-null   object  
 5   Cpu              1303 non-null   object  
 6   Ram              1303 non-null   object  
 7   Memory           1303 non-null   object  
 8   Gpu              1303 non-null   object  
 9   OpSys            1303 non-null   object  
 10  Weight           1303 non-null   object  
 11  Price             1303 non-null   float64
dtypes: float64(2), int64(1), object(9)
memory usage: 122.3+ KB
```

In [574]:

```
df.describe()
```

Out[574]:

| | Unnamed: 0 | Inches | Price |
|-------|-------------|-------------|--------------|
| count | 1303.000000 | 1303.000000 | 1303.000000 |
| mean | 651.000000 | 15.017191 | 59870.042910 |
| std | 376.28801 | 1.426304 | 37243.201786 |
| min | 0.00000 | 10.100000 | 9270.720000 |
| 25% | 325.50000 | 14.000000 | 31914.720000 |
| 50% | 651.00000 | 15.600000 | 52054.560000 |
| 75% | 976.50000 | 15.600000 | 79274.246400 |

```
max Unnamed: 0    Inches  ScreenResolution  Price
```

In [575]:

```
df.isnull().sum()
```

Out[575]:

```
Unnamed: 0      0  
Company        0  
TypeName       0  
Inches         0  
ScreenResolution 0  
Cpu            0  
Ram            0  
Memory         0  
Gpu            0  
OpSys          0  
Weight          0  
Price           0  
dtype: int64
```

In []:

In []:

In []:

DATA CLEANING:

In [576]:

```
df.head()
```

Out[576]:

| Unnamed: 0 | Company | TypeName | Inches | ScreenResolution | Cpu | Ram | Memory | Gpu | OpSys | Weight | Price | |
|---------------|---------|----------|-----------|------------------|---------------------------------------|----------------------------|--------|---------------------|------------------------------|--------|--------|-----------|
| 0 | 0 | Apple | Ultrabook | 13.3 | IPS Panel Retina Display 2560x1600 | Intel Core i5 2.3GHz | 8GB | 128GB SSD | Intel Iris Plus Graphics 640 | macOS | 1.37kg | 71378.68 |
| 1 | 1 | Apple | Ultrabook | 13.3 | 1440x900 | Intel Core i5 1.8GHz | 8GB | 128GB Flash Storage | Intel HD Graphics 6000 | macOS | 1.34kg | 47895.52 |
| 2 | 2 | HP | Notebook | 15.6 | Full HD 1920x1080 | Intel Core i5 7200U 2.5GHz | 8GB | 256GB SSD | Intel HD Graphics 620 | No OS | 1.86kg | 30636.00 |
| 3 | 3 | Apple | Ultrabook | 15.4 | IPS Panel Retina Display 2880x1800 | Intel Core i7 2.7GHz | 16GB | 512GB SSD | AMD Radeon Pro 455 | macOS | 1.83kg | 135195.33 |
| 4 | 4 | Apple | Ultrabook | 13.3 | IPS Panel Retina Display 2560x1600 | Intel Core i5 3.1GHz | 8GB | 256GB SSD | Intel Iris Plus Graphics 650 | macOS | 1.37kg | 96095.80 |

In []:

In [577]:

```
df.drop(columns=["Unnamed: 0"], inplace=True)
df["Ram"] = df["Ram"].str.replace("GB", "")
df["Weight"] = df["Weight"].str.replace("kg", "")

df["Ram"] = df["Ram"].astype(int)
df["Weight"] = df["Weight"].astype(float)
df["Inches"] = df["Inches"].astype(float)
df["Price"] = df["Price"].astype(float)
```

In [578]:

```
df.head()
```

Out [578]:

| | Company | TypeName | Inches | ScreenResolution | Cpu | Ram | Memory | Gpu | OpSys | Weight | Price |
|---|---------|-----------|--------|------------------------------------|----------------------------|-----|---------------------|------------------------------|-------|--------|-------------|
| 0 | Apple | Ultrabook | 13.3 | IPS Panel Retina Display 2560x1600 | Intel Core i5 2.3GHz | 8 | 128GB SSD | Intel Iris Plus Graphics 640 | macOS | 1.37 | 71378.6832 |
| 1 | Apple | Ultrabook | 13.3 | 1440x900 | Intel Core i5 1.8GHz | 8 | 128GB Flash Storage | Intel HD Graphics 6000 | macOS | 1.34 | 47895.5232 |
| 2 | HP | Notebook | 15.6 | Full HD 1920x1080 | Intel Core i5 7200U 2.5GHz | 8 | 256GB SSD | Intel HD Graphics 620 | No OS | 1.86 | 30636.0000 |
| 3 | Apple | Ultrabook | 15.4 | IPS Panel Retina Display 2880x1800 | Intel Core i7 2.7GHz | 16 | 512GB SSD | AMD Radeon Pro 455 | macOS | 1.83 | 135195.3360 |
| 4 | Apple | Ultrabook | 13.3 | IPS Panel Retina Display 2560x1600 | Intel Core i5 3.1GHz | 8 | 256GB SSD | Intel Iris Plus Graphics 650 | macOS | 1.37 | 96095.8080 |

In [579]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1303 entries, 0 to 1302
Data columns (total 11 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Company          1303 non-null   object 
 1   TypeName         1303 non-null   object 
 2   Inches           1303 non-null   float64
 3   ScreenResolution 1303 non-null   object 
 4   Cpu              1303 non-null   object 
 5   Ram              1303 non-null   int64  
 6   Memory           1303 non-null   object 
 7   Gpu              1303 non-null   object 
 8   OpSys            1303 non-null   object 
 9   Weight            1303 non-null   float64
 10  Price             1303 non-null   float64
dtypes: float64(3), int64(1), object(7)
memory usage: 112.1+ KB
```

In []:

In []:

EDA , FEATURE ENGINEERING AND DATA TRANSFORMATION

Section:

In [580]:

```
plt.figure(figsize=(25, 20))

data1 = df["Company"].value_counts().reset_index()

data2 = df["OpSys"].value_counts().reset_index()

plt.suptitle(
    "Exploratory Data Analysis of Laptop Dataset (Part - 1)",
    fontsize=24,
    fontweight="bold",
    y=0.98
)
plt.subplot(2,2,1)
sns.barplot(x="count",y="Company",data=data1.reset_index())
plt.xlabel("Production")
plt.ylabel("Company name")
plt.title("Company-wise Production")
```

#Each operating-system's contribution:


```
plt.subplot(2,2,2)
sns.barplot(x="count",y="OpSys",data=data2.reset_index())
plt.xlabel("Contribution")
plt.ylabel("Operating System")
plt.title("Operating System wise contribution")
```





```
plt.subplot(2,2,3)
data = df["TypeName"].value_counts()
plt.pie(x=data,labels=data.index,autopct=lambda p: f'{int(p*sum(data)/100)}',
        startangle=90)
plt.title("Contribution of each type of Laptops")
```

```
plt.subplot(2,2,4)
sns.distplot(df["Price"])
plt.title("Price Distribution of all laptops")
```

```
plt.show()
```

Exploratory Data Analysis of Laptop Dataset (Part - 1)





In [581]:

```

plt.figure(figsize=(25,20))
plt.suptitle(
    "Exploratory Data Analysis of Laptop Dataset (Part - 2)",
    fontsize=24,
    fontweight="bold",
    y=0.98
)

plt.tight_layout()
plt.subplot(2,2,1)
data = df.groupby("Company") ["Price"].mean().reset_index().sort_values(by="Price", ascending=False)
sns.barplot(y="Company",x="Price",data=data)
plt.title("Average Price Value of each company")
plt.xlabel("Average Price")
plt.ylabel("Company")

plt.subplot(2,2,2)
data = df[ "TypeName" ].value_counts().reset_index()
sns.barplot(y="TypeName",x="count",data=data)
plt.xlabel("Contribution")
plt.ylabel("Type of Laptop")


plt.subplot(2,2,3)
data = df.groupby("TypeName") ["Price"].mean().reset_index().sort_values(by="Price", ascending=False)
sns.barplot(y="TypeName",x="Price",data=data)
plt.title("Average Price Value of each laptop-type")
plt.xlabel("Average Price")
plt.ylabel("Laptop Type")
plt.tight_layout()

plt.subplot(2,2,4)
sns.lineplot(x=df[ "Inches" ],y=df[ "Price" ])
plt.title("Relationship between Laptop's screen size and its price")
plt.xlabel("Size (in inches)")
plt.ylabel("Laptop's price")
plt.tight_layout()
plt.show()

```

Exploratory Data Analysis of Laptop Dataset (Part - 2)





In []:

In []:

Now we will determine how many laptops are touch screen and how many arent :

(for this we will do feature engineering (i.e. create a new column named as: 'Touchscreen'))

In [582]:

```
df["ScreenResolution"].value_counts()
```

Out[582]:

| | |
|---|-----|
| ScreenResolution | |
| Full HD 1920x1080 | 507 |
| 1366x768 | 281 |
| IPS Panel Full HD 1920x1080 | 230 |
| IPS Panel Full HD / Touchscreen 1920x1080 | 53 |
| Full HD / Touchscreen 1920x1080 | 47 |
| 1600x900 | 23 |
| Touchscreen 1366x768 | 16 |
| Quad HD+ / Touchscreen 3200x1800 | 15 |
| IPS Panel 4K Ultra HD 3840x2160 | 12 |
| IPS Panel 4K Ultra HD / Touchscreen 3840x2160 | 11 |
| 4K Ultra HD / Touchscreen 3840x2160 | 10 |
| IPS Panel 1366x768 | 7 |
| Touchscreen 2560x1440 | 7 |
| 4K Ultra HD 3840x2160 | 7 |
| IPS Panel Retina Display 2304x1440 | 6 |

```

IPS Panel Retina Display 2560x1600          6
Touchscreen 2256x1504                      6
IPS Panel Quad HD+ / Touchscreen 3200x1800  6
IPS Panel Touchscreen 2560x1440              5
IPS Panel Retina Display 2880x1800           4
1440x900                                     4
IPS Panel Touchscreen 1920x1200              4
IPS Panel 2560x1440                          4
IPS Panel Quad HD+ 2560x1440                 3
IPS Panel Touchscreen 1366x768               3
Quad HD+ 3200x1800                          3
1920x1080                                    3
2560x1440                                    3
Touchscreen 2400x1600                        3
IPS Panel Quad HD+ 3200x1800                 2
IPS Panel Full HD 2160x1440                  2
IPS Panel Touchscreen / 4K Ultra HD 3840x2160 2
IPS Panel Full HD 1366x768                  1
Touchscreen / Quad HD+ 3200x1800             1
IPS Panel Retina Display 2736x1824            1
IPS Panel Full HD 2560x1440                  1
IPS Panel Full HD 1920x1200                  1
Touchscreen / Full HD 1920x1080              1
Touchscreen / 4K Ultra HD 3840x2160            1
IPS Panel Touchscreen 2400x1600              1
Name: count, dtype: int64

```

In [583]:

```
# 0 if not touchscreen and 1 if touchscreen:
df["Touchscreen"] = df["ScreenResolution"].apply(lambda x: 1 if 'Touchscreen' in x else 0)
```

In [584]:

```
#inspect first 5 rows of the dataframe:
df.head()
```

Out[584]:

| | Company | Type | Name | Inches | ScreenResolution | Cpu | Ram | Memory | Gpu | OpSys | Weight | Price | Touchscr |
|---|---------|-----------|------|--------|------------------------------------|----------------------------|-----|---------------------|------------------------------|-------|--------|-------------|----------|
| 0 | Apple | Ultrabook | 13.3 | | IPS Panel Retina Display 2560x1600 | Intel Core i5 2.3GHz | 8 | 128GB SSD | Intel Iris Plus Graphics 640 | macOS | 1.37 | 71378.6832 | |
| 1 | Apple | Ultrabook | 13.3 | | 1440x900 | Intel Core i5 1.8GHz | 8 | 128GB Flash Storage | Intel HD Graphics 6000 | macOS | 1.34 | 47895.5232 | |
| 2 | HP | Notebook | 15.6 | | Full HD 1920x1080 | Intel Core i5 7200U 2.5GHz | 8 | 256GB SSD | Intel HD Graphics 620 | No OS | 1.86 | 30636.0000 | |
| 3 | Apple | Ultrabook | 15.4 | | IPS Panel Retina Display 2880x1800 | Intel Core i7 2.7GHz | 16 | 512GB SSD | AMD Radeon Pro 455 | macOS | 1.83 | 135195.3360 | |
| 4 | Apple | Ultrabook | 13.3 | | IPS Panel Retina Display 2560x1600 | Intel Core i5 3.1GHz | 8 | 256GB SSD | Intel Iris Plus Graphics 650 | macOS | 1.37 | 96095.8080 | |

In [585]:

```
#take a sample of any 5 rows:
df.sample(5)
```

Out[585]:

| Company | Type | Name | Inches | Screen Resolution | Cpu | Ram | Memory | Gpu | OpSys | Weight | Price | Touchscreen |
|---------|--------|----------|--------|--------------------------------|-----------------------------------|-----|---------------------|-------------------------|------------|--------|----------|-------------|
| Company | Type | Name | Inches | Screen Resolution | Cpu | Ram | Memory | Gpu | OpSys | Weight | Price | Touchscreen |
| 31 | Asus | Notebook | 14.0 | 1366x768 | AMD E-Series E2-6110 1.5GHz | 2 | 32GB Flash Storage | AMD Radeon R2 | Windows 10 | 1.65 | 10602.72 | |
| 557 | Lenovo | Notebook | 17.3 | 1600x900 | Intel Core i7 7500U 2.7GHz | 6 | 128GB SSD + 1TB HDD | Nvidia GeForce 940MX | Windows 10 | 2.80 | 50562.72 | |
| 842 | HP | Notebook | 17.3 | 1600x900 | Intel Core i5 7200U 2.5GHz | 8 | 1TB HDD | Nvidia GeForce 930MX | Windows 10 | 2.63 | 48484.80 | |
| 105 | HP | Notebook | 14.0 | IPS Panel Full HD 1920x1080 | Intel Core i5 7200U 2.5GHz | 6 | 256GB SSD | Nvidia GeForce 940MX | Windows 10 | 1.58 | 35111.52 | |
| 75 | Asus | Gaming | 15.6 | Full HD 1920x1080 | Intel Core i7 7700HQ 2.8GHz | 8 | 1TB HDD | Nvidia GeForce GTX 1050 | Windows 10 | 2.20 | 50562.72 | |

[1]

In []:

In [586]:

```
df[ "Touchscreen" ].value_counts().reset_index()
```

Out[586]:

| Touchscreen count | | |
|-------------------|---|------|
| 0 | 0 | 1111 |
| 1 | 1 | 192 |

In [587]:

```
plt.figure(figsize=(20,15))

plt.suptitle(
    "Exploratory Data Analysis of Laptop Dataset (Part - 3)",
    fontsize=24,
    fontweight="bold",
    y=0.98
)

plt.subplot(2,2,1)
data = df[ "Touchscreen" ].value_counts().reset_index()

labels = data[ "Touchscreen" ].map({
    0: "Non-Touchscreen Laptops",
    1: "Touchscreen Laptops"
})

plt.pie(
    data[ "count" ],
    labels=labels,
    autopct=lambda p: f'{int(p*sum(data[ "count" ])/100)}',
    startangle=90
)
```

```

plt.title("Touchscreen vs Non-Touchscreen Laptops")

plt.subplot(2,2,2)
plt.title("Frequency of both Touchscreen and non-touchscreen laptops")
sns.barplot(x="Touchscreen", y="count", data=data)





for c in plt.gca().containers:
    plt.gca().bar_label(c, padding=3)

plt.subplot(2,2,3)
plt.title("Average Price of both Touchscreen and non-touchscreen laptops")
### Now we will see average prices of both touch-screen and non-touchscreen laptops:_
data = df.groupby("Touchscreen")["Price"].mean().reset_index(name = "Average Price")
sns.barplot(y="Average Price",x="Touchscreen",data=data)
for c in plt.gca().containers:
    plt.gca().bar_label(c, padding=3)

plt.subplot(2,2,4)
plt.title("Laptops with and without IPS screen resolution")
df["Ips"] = df["ScreenResolution"].apply(lambda x:1 if 'IPS' in x else 0)
data = df["Ips"].value_counts().reset_index()

labels = data["Ips"].map({
    0: "Laptops without IPS Screen Resolution",
    1: "Laptops with IPS Screen Resolution"
})
plt.pie(
    data["count"],
    labels=labels,
    autopct=lambda p: f'{int(p*sum(data["count"])/100)}',
    startangle=90
)
plt.show()

```



In []:

Extract screen-resolution's size (means x and y (in inches)) from 'ScreenResolution' column:

In [588]:

```
df.head()
```

Out [588]:

| Company | Type | Name | Inches | ScreenResolution | Cpu | Ram | Memory | Gpu | OpSys | Weight | Price | Touchscr |
|---------|-------|-----------|--------|---------------------------------------|----------------------------|-----|---------------------|------------------------------|-------|--------|-------------|----------|
| 0 | Apple | Ultrabook | 13.3 | IPS Panel Retina Display 2560x1600 | Intel Core i5 2.3GHz | 8 | 128GB SSD | Intel Iris Plus Graphics 640 | macOS | 1.37 | 71378.6832 | |
| 1 | Apple | Ultrabook | 13.3 | 1440x900 | Intel Core i5 1.8GHz | 8 | 128GB Flash Storage | Intel HD Graphics 6000 | macOS | 1.34 | 47895.5232 | |
| 2 | HP | Notebook | 15.6 | Full HD 1920x1080 | Intel Core i5 7200U 2.5GHz | 8 | 256GB SSD | Intel HD Graphics 620 | No OS | 1.86 | 30636.0000 | |
| 3 | Apple | Ultrabook | 15.4 | IPS Panel Retina Display 2880x1800 | Intel Core i7 2.7GHz | 16 | 512GB SSD | AMD Radeon Pro 455 | macOS | 1.83 | 135195.3360 | |
| 4 | Apple | Ultrabook | 13.3 | IPS Panel Retina Display | Intel Core i5 | 8 | 256GB SSD | Intel Iris Plus Graphics | macOS | 1.37 | 96095.8080 | |

| Company | TypeName | Inches | ScreenResolution | 2560x1600 | 3.1GHz | Cpu | Ram | Memory | 650 | OpSys | Weight | Price | Touchscr |
|---------|----------|--------|------------------|-----------|--------|-----|-----|--------|-----|-------|--------|-------|----------|
|---------|----------|--------|------------------|-----------|--------|-----|-----|--------|-----|-------|--------|-------|----------|

In [589]:

```
new = df["ScreenResolution"].str.split('x', n=1, expand=True)
```

In [590]:

```
df["X_res"] = new[0]
df["Y_res"] = new[1]
```

In [591]:

```
df.head()
```

Out[591]:

| Company | TypeName | Inches | ScreenResolution | Cpu | Ram | Memory | Gpu | OpSys | Weight | Price | Touchscr |
|---------|----------|-----------|------------------|------------------------------------|----------------------------|--------|---------------------|------------------------------|--------|-------|-------------|
| 0 | Apple | Ultrabook | 13.3 | IPS Panel Retina Display 2560x1600 | Intel Core i5 2.3GHz | 8 | 128GB SSD | Intel Iris Plus Graphics 640 | macOS | 1.37 | 71378.6832 |
| 1 | Apple | Ultrabook | 13.3 | 1440x900 | Intel Core i5 1.8GHz | 8 | 128GB Flash Storage | Intel HD Graphics 6000 | macOS | 1.34 | 47895.5232 |
| 2 | HP | Notebook | 15.6 | Full HD 1920x1080 | Intel Core i5 7200U 2.5GHz | 8 | 256GB SSD | Intel HD Graphics 620 | No OS | 1.86 | 30636.0000 |
| 3 | Apple | Ultrabook | 15.4 | IPS Panel Retina Display 2880x1800 | Intel Core i7 2.7GHz | 16 | 512GB SSD | AMD Radeon Pro 455 | macOS | 1.83 | 135195.3360 |
| 4 | Apple | Ultrabook | 13.3 | IPS Panel Retina Display 2560x1600 | Intel Core i5 3.1GHz | 8 | 256GB SSD | Intel Iris Plus Graphics 650 | macOS | 1.37 | 96095.8080 |

In [592]:

```
df["X_res"] = df["X_res"].str.replace(", ", "").str.findall(r'(\d+\.\?\d+)').apply(lambda x : x[0])
```

In [593]:

```
df.head()
```

Out[593]:

| Company | TypeName | Inches | ScreenResolution | Cpu | Ram | Memory | Gpu | OpSys | Weight | Price | Touchscr |
|---------|----------|-----------|------------------|------------------------------------|----------------------------|--------|---------------------|------------------------------|--------|-------|------------|
| 0 | Apple | Ultrabook | 13.3 | IPS Panel Retina Display 2560x1600 | Intel Core i5 2.3GHz | 8 | 128GB SSD | Intel Iris Plus Graphics 640 | macOS | 1.37 | 71378.6832 |
| 1 | Apple | Ultrabook | 13.3 | 1440x900 | Intel Core i5 1.8GHz | 8 | 128GB Flash Storage | Intel HD Graphics 6000 | macOS | 1.34 | 47895.5232 |
| 2 | HP | Notebook | 15.6 | Full HD 1920x1080 | Intel Core i5 7200U 2.5GHz | 8 | 256GB SSD | Intel HD Graphics 620 | No OS | 1.86 | 30636.0000 |

| Company | Type Name | Inches | Screen Resolution | Gpu | Ram | Memory | Brand | OpSys | Weight | Price | Touchscr |
|---------|-----------|-----------|-------------------|--|----------------------------|--------|--------------|---------------------------------------|--------|-------|-------------|
| 3 | Apple | Ultrabook | 15.4 | Display 2880x1800 | Core i7 2.7GHz | 16 | 512GB | Radeon | macOS | 1.83 | 135195.3360 |
| 4 | Apple | Ultrabook | 13.3 | IPS Panel Retina Display 2560x1600 | Intel Core i5 3.1GHz | 8 | 256GB SSD | Intel Iris Plus Graphics 650 | macOS | 1.37 | 96095.8080 |

In [594]:

```
df["X_res"] = df["X_res"].apply(int)
df["Y_res"] = df["Y_res"].apply(int)
```

In [595]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1303 entries, 0 to 1302
Data columns (total 15 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Company          1303 non-null   object 
 1   TypeName         1303 non-null   object 
 2   Inches           1303 non-null   float64
 3   ScreenResolution 1303 non-null   object 
 4   Cpu              1303 non-null   object 
 5   Ram              1303 non-null   int64  
 6   Memory           1303 non-null   object 
 7   Gpu              1303 non-null   object 
 8   OpSys            1303 non-null   object 
 9   Weight            1303 non-null   float64
 10  Price             1303 non-null   float64
 11  Touchscreen      1303 non-null   int64  
 12  Ips              1303 non-null   int64  
 13  X_res            1303 non-null   int64  
 14  Y_res            1303 non-null   int64  
dtypes: float64(3), int64(5), object(7)
memory usage: 152.8+ KB
```

In [596]:

```
#Correlation of all numeric columns with respect to 'Price' column:
df.corr(numeric_only=True) ['Price']
```

Out[596]:

```
Inches          0.068197
Ram            0.743007
Weight          0.210370
Price           1.000000
Touchscreen    0.191226
Ips             0.252208
X_res           0.556529
Y_res           0.552809
Name: Price, dtype: float64
```

Prepare a new-column named 'PPI' (stands for Pixels-Per-Inches):

In [597]:

```
df["ppi"] = (((df['X_res']**2) + (df['Y_res']**2))**0.5 / df['Inches']).astype('float')
```

In [598]:

```
df.head()
```

Out[598]:

| Company | Type Name | Inches | Screen Resolution | Cpu | Ram | Memory | Gpu | OpSys | Weight | Price | Touchscr |
|---------|-----------|-----------|-------------------|---------------------------------------|----------------------------|-----------------------|------------------------------|-------|--------|-------------|----------|
| 0 | Apple | Ultrabook | 13.3 | IPS Panel Retina Display 2560x1600 | Intel Core i5 2.3GHz | 8 128GB SSD | Intel Iris Plus Graphics 640 | macOS | 1.37 | 71378.6832 | |
| 1 | Apple | Ultrabook | 13.3 | 1440x900 | Intel Core i5 1.8GHz | 8 128GB Flash Storage | Intel HD Graphics 6000 | macOS | 1.34 | 47895.5232 | |
| 2 | HP | Notebook | 15.6 | Full HD 1920x1080 | Intel Core i5 7200U 2.5GHz | 8 256GB SSD | Intel HD Graphics 620 | No OS | 1.86 | 30636.0000 | |
| 3 | Apple | Ultrabook | 15.4 | IPS Panel Retina Display 2880x1800 | Intel Core i7 2.7GHz | 16 512GB SSD | AMD Radeon Pro 455 | macOS | 1.83 | 135195.3360 | |
| 4 | Apple | Ultrabook | 13.3 | IPS Panel Retina Display 2560x1600 | Intel Core i5 3.1GHz | 8 256GB SSD | Intel Iris Plus Graphics 650 | macOS | 1.37 | 96095.8080 | |

In [599]:

```
# Correlation of all numeric columns with respect to 'Price' column:
df.corr(numeric_only=True) ['Price']
```

Out[599]:

```
Inches          0.068197
Ram            0.743007
Weight         0.210370
Price          1.000000
Touchscreen    0.191226
Ips            0.252208
X_res          0.556529
Y_res          0.552809
ppi            0.473487
Name: Price, dtype: float64
```

In [600]:

```
# Drop 'ScreenResolution' column:
df.drop(columns='ScreenResolution', inplace=True)
```

In [601]:

```
df.columns
```

Out[601]:

```
Index(['Company', 'Type Name', 'Inches', 'Cpu', 'Ram', 'Memory', 'Gpu', 'OpSys',
       'Weight', 'Price', 'Touchscreen', 'Ips', 'X_res', 'Y_res', 'ppi'],
      dtype='object')
```

In [602]:

```
# Drop 'Inches', 'X_res', 'Y_res' columns:
df.drop(columns=["Inches", "X_res", "Y_res"], inplace=True)
```

In [603]:

```
df.columns
```

Out[603]:

```
Index(['Company', 'Type Name', 'Cpu', 'Ram', 'Memory', 'Gpu', 'OpSys', 'Weight',
       'Price', 'Touchscreen', 'Ips', 'ppi'],
      dtype='object')
```

In []:

In [604]:

```
#simplify the 'Cpu' column:
```

```
#step 1) extract first 3 words of all the values:
```

```
df["Cpu Name"] = df["Cpu"].apply(lambda x:" ".join(x.split()[0:3]))
```

In [605]:

```
df.head()
```

Out[605]:

| | Company | TypeName | Cpu | Ram | Memory | Gpu | OpSys | Weight | Price | Touchscreen | Ips | ppi | Cp Name |
|---|---------|-----------|----------------------------|-----|---------------------|------------------------------|-------|--------|-------------|-------------|-----|------------|--------------|
| 0 | Apple | Ultrabook | Intel Core i5 2.3GHz | 8 | 128GB SSD | Intel Iris Plus Graphics 640 | macOS | 1.37 | 71378.6832 | 0 | 1 | 226.983005 | Intel Core i |
| 1 | Apple | Ultrabook | Intel Core i5 1.8GHz | 8 | 128GB Flash Storage | Intel HD Graphics 6000 | macOS | 1.34 | 47895.5232 | 0 | 0 | 127.677940 | Intel Core i |
| 2 | HP | Notebook | Intel Core i5 7200U 2.5GHz | 8 | 256GB SSD | Intel HD Graphics 620 | No OS | 1.86 | 30636.0000 | 0 | 0 | 141.211998 | Intel Core i |
| 3 | Apple | Ultrabook | Intel Core i7 2.7GHz | 16 | 512GB SSD | AMD Radeon Pro 455 | macOS | 1.83 | 135195.3360 | 0 | 1 | 220.534624 | Intel Core i |
| 4 | Apple | Ultrabook | Intel Core i5 3.1GHz | 8 | 256GB SSD | Intel Iris Plus Graphics 650 | macOS | 1.37 | 96095.8080 | 0 | 1 | 226.983005 | Intel Core i |

In [606]:

```
def fetch_processor(text):
    if text == 'Intel Core i7' or text == 'Intel Core i5' or text == 'Intel Core i3':
        return text
    else:
        if text.split()[0] == 'Intel':
            return 'Other Intel Processor'
        else:
            return 'AMD Processor'
```

In [607]:

```
df["Cpu brand"] = df["Cpu Name"].apply(fetch_processor)
```

In [608]:

```
df.head()
```

Out[608]:

| | Company | TypeName | Cpu | Ram | Memory | Gpu | OpSys | Weight | Price | Touchscreen | Ips | ppi | Cp Name |
|---|---------|-----------|----------------------|-----|-----------|------------------------------|-------|--------|------------|-------------|-----|------------|--------------|
| 0 | Apple | Ultrabook | Intel Core i5 2.3GHz | 8 | 128GB SSD | Intel Iris Plus Graphics 640 | macOS | 1.37 | 71378.6832 | 0 | 1 | 226.983005 | Intel Core i |
| | | | Intel | | 128GB | Intel HD | | | | | | | Intel |

| 1 | Apple Company | Ultrabook Type | Name | Core i5 1.8GHz | Ram 8 Gb | Flash Storage 256GB SSD | Graphics 620 | macOS OpSys | Weight 1.34 | Price 47895.5232 | Touchscreen 0 | Lips 0 | 127.677940 ppi | Op Name |
|---|---------------|----------------|----------------------------|----------------|-----------|------------------------------|--------------|-------------|-------------|------------------|---------------|------------|----------------|---------|
| 2 | HP | Notebook | Intel Core i5 7200U 2.5GHz | 8 | 256GB SSD | Intel HD Graphics 620 | No OS | 1.86 | 30636.0000 | 0 | 0 | 141.211998 | Inte Cor i | |
| 3 | Apple | Ultrabook | Intel Core i7 2.7GHz | 16 | 512GB SSD | AMD Radeon Pro 455 | macOS | 1.83 | 135195.3360 | 0 | 1 | 220.534624 | Inte Cor i | |
| 4 | Apple | Ultrabook | Intel Core i5 3.1GHz | 8 | 256GB SSD | Intel Iris Plus Graphics 650 | macOS | 1.37 | 96095.8080 | 0 | 1 | 226.983005 | Inte Cor i | |

In [609]:

```
plt.figure(figsize=(20, 7))


plt.suptitle(
    "Exploratory Data Analysis of Laptop Dataset (Part - 4)",
    fontsize=24,
    fontweight="bold",
    y=0.98
)

plt.subplot(2, 2, 1)
sns.countplot(df["Cpu brand"])

plt.subplot(2, 2, 3)
sns.barplot(y="Cpu brand", x="Price", data=df)

plt.show()
```

Exploratory Data Analysis of Laptop Dataset (Part - 4)



In []:

```
# Drop 'Cpu' and 'Cpu Name' columns (as they aren't required now):
df.drop(columns=["Cpu", "Cpu Name"], inplace=True)
```

In [611]:

```
df.head()
```

Out[611]:

| | Company | Type | Name | Ram | Memory | Gpu | OpSys | Weight | Price | Touchscreen | Ips | ppi | Cpu brand |
|---|---------|-----------|------|-----|---------------------|------------------------------|-------|--------|-------------|-------------|-----|------------|---------------|
| 0 | Apple | Ultrabook | | 8 | 128GB SSD | Intel Iris Plus Graphics 640 | macOS | 1.37 | 71378.6832 | 0 | 1 | 226.983005 | Intel Core i5 |
| 1 | Apple | Ultrabook | | 8 | 128GB Flash Storage | Intel HD Graphics 6000 | macOS | 1.34 | 47895.5232 | 0 | 0 | 127.677940 | Intel Core i5 |
| 2 | HP | Notebook | | 8 | 256GB SSD | Intel HD Graphics 620 | No OS | 1.86 | 30636.0000 | 0 | 0 | 141.211998 | Intel Core i5 |
| 3 | Apple | Ultrabook | | 16 | 512GB SSD | AMD Radeon Pro 455 | macOS | 1.83 | 135195.3360 | 0 | 1 | 220.534624 | Intel Core i7 |
| 4 | Apple | Ultrabook | | 8 | 256GB SSD | Intel Iris Plus Graphics 650 | macOS | 1.37 | 96095.8080 | 0 | 1 | 226.983005 | Intel Core i5 |

In []:

In [612]:

```
plt.figure(figsize=(15,10))
```


```
plt.suptitle(
    "Exploratory Data Analysis of Laptop Dataset (Part - 5)",
    fontsize=24,
    fontweight="bold",
    y=0.98
)
```

```
# Which Ram is mostly soled in the market:
plt.subplot(2,2,1)
order = df['Ram'].value_counts().index
plt.title("Most sold RAM")
sns.countplot(x="Ram", data=df, order=order)
plt.xlabel("RAM Size")
plt.ylabel("Quantity sold")
```



```
# Which Ram is most costly?
plt.subplot(2,2,2)
sns.barplot(x="Ram", y="Price", data=df)
plt.title("RAM-size and it's cost")
plt.xlabel("RAM Size")
plt.ylabel("Cost")
```

```
# Trend between RAM-size and it's cost:
plt.subplot(2,2,3)
sns.lineplot(x="Ram", y="Price", data=df)
plt.title("Trend between RAM-size and it's cost")
plt.xlabel("RAM Size")
plt.ylabel("Cost")
plt.show()
```

Exploratory Data Analysis of Laptop Dataset (Part - 5)



Trend between RAM-size and it's cost



In []:

In []:

In [613]:

df.head()

Out[613]:

| | Company | TypeName | Ram | Memory | Gpu | OpSys | Weight | Price | Touchscreen | Ips | ppi | Cpu brand |
|---|---------|-----------|-----|---------------------|------------------------------|-------|--------|-------------|-------------|-----|------------|---------------|
| 0 | Apple | Ultrabook | 8 | 128GB SSD | Intel Iris Plus Graphics 640 | macOS | 1.37 | 71378.6832 | 0 | 1 | 226.983005 | Intel Core i5 |
| 1 | Apple | Ultrabook | 8 | 128GB Flash Storage | Intel HD Graphics 6000 | macOS | 1.34 | 47895.5232 | 0 | 0 | 127.677940 | Intel Core i5 |
| 2 | HP | Notebook | 8 | 256GB SSD | Intel HD Graphics 620 | No OS | 1.86 | 30636.0000 | 0 | 0 | 141.211998 | Intel Core i5 |
| 3 | Apple | Ultrabook | 16 | 512GB SSD | AMD Radeon Pro 455 | macOS | 1.83 | 135195.3360 | 0 | 1 | 220.534624 | Intel Core i7 |
| 4 | Apple | Ultrabook | 8 | 256GB SSD | Intel Iris Plus Graphics 650 | macOS | 1.37 | 96095.8080 | 0 | 1 | 226.983005 | Intel Core i5 |

In [614]:

df["Memory"].value_counts()

Out[614]:

```
Memory
256GB SSD          412
1TB HDD            223
500GB HDD          132
512GB SSD          118
128GB SSD + 1TB HDD 94
128GB SSD          76
256GB SSD + 1TB HDD 73
32GB Flash Storage 38
2TB HDD            16
64GB Flash Storage 15
1TB SSD            14
512GB SSD + 1TB HDD 14
256GB SSD + 2TB HDD 10
1.0TB Hybrid        9
256GB Flash Storage 8
16GB Flash Storage 7
32GB SSD            6
180GB SSD           5
128GB Flash Storage 4
16GB SSD            3
512GB SSD + 2TB HDD 3
128GB SSD + 2TB HDD 2
256GB SSD + 256GB SSD 2
512GB Flash Storage 2
1TB SSD + 1TB HDD   2
256GB SSD + 500GB HDD 2
64GB SSD            1
512GB SSD + 512GB SSD 1
64GB Flash Storage + 1TB HDD 1
1TB HDD + 1TB HDD   1
512GB SSD + 256GB SSD 1
32GB HDD            1
128GB HDD           1
240GB SSD           1
8GB SSD             1
508GB Hybrid         1
1.0TB HDD           1
512GB SSD + 1.0TB Hybrid 1
256GB SSD + 1.0TB Hybrid 1
Name: count, dtype: int64
```

NOW WE WILL DO EDA ON 'Memory' COLUMN:

EARLIER:

In [615]:

```
df.head(1)
```

Out[615]:

| | Company | TypeName | Ram | Memory | Gpu | OpSys | Weight | Price | Touchscreen | Ips | ppi | Cpu brand |
|---|---------|-----------|-----|-----------|------------------------------|-------|--------|------------|-------------|-----|------------|---------------|
| 0 | Apple | Ultrabook | 8 | 128GB SSD | Intel Iris Plus Graphics 640 | macOS | 1.37 | 71378.6832 | 0 | 1 | 226.983005 | Intel Core i5 |

NOW AFTER TRANSFORMATION:

In [616]:

```
df["Memory"] = df["Memory"].astype(str).replace('.0', '', regex=True)
df["Memory"] = df["Memory"].str.replace("GB", "")
```

```

df["Memory"] = df["Memory"].str.replace("TB", "000")

new = df["Memory"].str.split("+", n=1, expand=True)

df["first"] = new[0]
df["first"] = df["first"].str.strip()
df["second"] = new[1]

df["Layer1HDD"] = df["first"].apply(lambda x: 1 if "HDD" in x else 0)
df["Layer1SSD"] = df["first"].apply(lambda x: 1 if "SSD" in x else 0)
df["Layer1Hybrid"] = df["first"].apply(lambda x: 1 if "Hybrid" in x else 0)
df["Layer1Flash_Storage"] = df["first"].apply(lambda x: 1 if "Flash Storage" in x else 0)

df["first"] = df["first"].str.replace(r"\D", "")

df["second"].fillna(0, inplace=True)

df["Layer2HDD"] = df["second"].apply(lambda x: 1 if "HDD" in x else 0)
df["Layer2SSD"] = df["second"].apply(lambda x: 1 if "SSD" in x else 0)
df["Layer2Hybrid"] = df["second"].apply(lambda x: 1 if "Hybrid" in x else 0)
df["Layer2Flash_Storage"] = df["second"].apply(lambda x: 1 if "Flash Storage" in x else 0)

df["second"] = df["second"].str.replace(r"\D", "")

df["first"] = (
    df["first"]
    .str.extract(r"(\d+)", expand=False)
    .astype(int)
)

df["second"] = (
    df["second"]
    .str.extract(r"(\d+)", expand=False)
    .fillna(0)
    .astype(int)
)

df["HDD"] = (df["first"] * df["Layer1HDD"] + df["second"] * df["Layer2HDD"])
df["SSD"] = (df["first"] * df["Layer1SSD"] + df["second"] * df["Layer2SSD"])
df["Hybrid"] = (df["first"] * df["Layer1Hybrid"] + df["second"] * df["Layer2Hybrid"])
df["Flash_Storage"] = (
    df["first"] * df["Layer1Flash_Storage"] +
    df["second"] * df["Layer2Flash_Storage"]
)

df.drop(
    columns=[
        "first", "second",
        "Layer1HDD", "Layer1SSD", "Layer1Hybrid", "Layer1Flash_Storage",
        "Layer2HDD", "Layer2SSD", "Layer2Hybrid", "Layer2Flash_Storage"
    ],
    inplace=True
)

```

In [617]:

```
df.head()
```

Out[617]:

| Company | Type | Name | Ram | Memory | Gpu | OpSys | Weight | Price | Touchscreen | Ips | ppi | Cpu brand | HDD |
|---------|-------|-----------|-----|-----------|------------------------------|-------|--------|------------|-------------|-----|------------|---------------|-----|
| 0 | Apple | Ultrabook | 8 | 128 SSD | Intel Iris Plus Graphics 640 | macOS | 1.37 | 71378.6832 | 0 | 1 | 226.983005 | Intel Core i5 | 0 |
| 1 | Apple | Ultrabook | 8 | 128 Flash | Intel HD Graphics | macOS | 1.24 | 47805.5022 | 0 | 0 | 107.677010 | Intel Core | 0 |

| Company | Type | Name | Ram | Storage | Gpu | OpSys | Weight | Price | Touchscreen | Ips | ppi | Cpu brand | HDD | SSD | Hybrid |
|---------|-------|-----------|-----|---------|------------------------------|-------|--------|-------------|-------------|-----|------------|---------------|-----|-----|--------|
| 2 | HP | Notebook | 8 | 256 SSD | Graphics 620 | No OS | 1.86 | 30636.0000 | 0 | 0 | 141.211998 | Core i5 | 0 | 0 | 0 |
| 3 | Apple | Ultrabook | 16 | 512 SSD | AMD Radeon Pro 455 | macOS | 1.83 | 135195.3360 | 0 | 1 | 220.534624 | Intel Core i7 | 0 | 0 | 0 |
| 4 | Apple | Ultrabook | 8 | 256 SSD | Intel Iris Plus Graphics 650 | macOS | 1.37 | 96095.8080 | 0 | 1 | 226.983005 | Intel Core i5 | 0 | 0 | 0 |

In [618]:

```
#Drop the 'Memory' column since it isn't required now:
df.drop(columns=["Memory"], inplace=True)
```

In [619]:

```
df.head()
```

Out[619]:

| Company | Type | Name | Ram | Gpu | OpSys | Weight | Price | Touchscreen | Ips | ppi | Cpu brand | HDD | SSD | Hybrid |
|---------|-------|-----------|-----|------------------------------|-------|--------|-------------|-------------|-----|------------|---------------|-----|-----|--------|
| 0 | Apple | Ultrabook | 8 | Intel Iris Plus Graphics 640 | macOS | 1.37 | 71378.6832 | 0 | 1 | 226.983005 | Intel Core i5 | 0 | 128 | 0 |
| 1 | Apple | Ultrabook | 8 | Intel HD Graphics 6000 | macOS | 1.34 | 47895.5232 | 0 | 0 | 127.677940 | Intel Core i5 | 0 | 0 | 0 |
| 2 | HP | Notebook | 8 | Intel HD Graphics 620 | No OS | 1.86 | 30636.0000 | 0 | 0 | 141.211998 | Intel Core i5 | 0 | 256 | 0 |
| 3 | Apple | Ultrabook | 16 | AMD Radeon Pro 455 | macOS | 1.83 | 135195.3360 | 0 | 1 | 220.534624 | Intel Core i7 | 0 | 512 | 0 |
| 4 | Apple | Ultrabook | 8 | Intel Iris Plus Graphics 650 | macOS | 1.37 | 96095.8080 | 0 | 1 | 226.983005 | Intel Core i5 | 0 | 256 | 0 |

In [620]:

```
#See correlation of all columns with respect to the 'Price' column:
df.corr(numeric_only=True) ['Price']
```

Out[620]:

```
Ram          0.743007
Weight       0.210370
Price        1.000000
Touchscreen  0.191226
Ips          0.252208
ppi          0.473487
HDD          -0.096441
SSD          0.670799
Hybrid       0.007989
Flash_Storage -0.040511
Name: Price, dtype: float64
```

In [621]:

```
#Drop the columns 'Hybrid', 'Flash Storage' since they serve no purpose:
```

```
df.drop(columns=["Hybrid", "Flash Storage"], inplace=True)
```

Now do eda on 'Gpu' and 'OpSys' columns:

In [622]:

```
df["Gpu"].value_counts().reset_index()
```

Out[622]:

| | Gpu | count |
|-----|-------------------------|-------|
| 0 | Intel HD Graphics 620 | 281 |
| 1 | Intel HD Graphics 520 | 185 |
| 2 | Intel UHD Graphics 620 | 68 |
| 3 | Nvidia GeForce GTX 1050 | 66 |
| 4 | Nvidia GeForce GTX 1060 | 48 |
| ... | ... | ... |
| 105 | Nvidia Quadro M500M | 1 |
| 106 | AMD Radeon R7 M360 | 1 |
| 107 | Nvidia Quadro M3000M | 1 |
| 108 | Nvidia GeForce 960M | 1 |
| 109 | ARM Mali T860 MP4 | 1 |

110 rows × 2 columns

In [623]:

```
df[\"OpSys\"].value_counts().reset_index()
```

Out[623]:

| | OpSys | count |
|---|--------------|-------|
| 0 | Windows 10 | 1072 |
| 1 | No OS | 66 |
| 2 | Linux | 62 |
| 3 | Windows 7 | 45 |
| 4 | Chrome OS | 27 |
| 5 | macOS | 13 |
| 6 | Mac OS X | 8 |
| 7 | Windows 10 S | 8 |
| 8 | Android | 2 |

In [624]:

```
df.head()
```

Out[624]:

| Company | Type Name | Ram | Gpu | OpSys | Weight | Price | Touchscreen | Ips | ppi | Cpu brand | HDD | SSD |
|---------|-----------|-----------|-----|------------------------------|--------|-------|-------------|-----|-----|------------|---------------|-------|
| 0 | Apple | Ultrabook | 8 | Intel Iris Plus Graphics 640 | macOS | 1.37 | 71378.6832 | 0 | 1 | 226.983005 | Intel Core i5 | 0 128 |
| | | | | Intel HD | | | | | | Intel | | |

| 1 | Apple Company | Ultrabook Type | 8 Ram | Graphics GPU | macOS OpSys | 1.34 Weight | 47895.5232 Price | 0 Touchscreen | 0 Ips | 127.677940 ppi | Core i5 Cpu brand | 0 HDD | 0 SSD |
|---|---------------|----------------|-------|------------------------------|-------------|-------------|------------------|---------------|-------|----------------|-------------------|-------|-------|
| 2 | HP | Notebook | 8 | Intel HD Graphics 620 | No OS | 1.86 | 30636.0000 | 0 | 0 | 141.211998 | Intel Core i5 | 0 | 256 |
| 3 | Apple | Ultrabook | 16 | Radeon Pro 455 | macOS | 1.83 | 135195.3360 | 0 | 1 | 220.534624 | Intel Core i7 | 0 | 512 |
| 4 | Apple | Ultrabook | 8 | Intel Iris Plus Graphics 650 | macOS | 1.37 | 96095.8080 | 0 | 1 | 226.983005 | Intel Core i5 | 0 | 256 |

In [625]:

```
df["Gpu brand"] = df["Gpu"].apply(lambda x:x.split()[0])
```

In [626]:

```
df.head()
```

Out[626]:

| Company | Type Name | Ram | Gpu | OpSys | Weight | Price | Touchscreen | Ips | ppi | Cpu brand | HDD | SSD | G bra | |
|---------|-----------|-----------|-----|------------------------------|--------|-------|-------------|-----|-----|------------|---------------|-----|-------|----|
| 0 | Apple | Ultrabook | 8 | Intel Iris Plus Graphics 640 | macOS | 1.37 | 71378.6832 | 0 | 1 | 226.983005 | Intel Core i5 | 0 | 128 | In |
| 1 | Apple | Ultrabook | 8 | Intel HD Graphics 6000 | macOS | 1.34 | 47895.5232 | 0 | 0 | 127.677940 | Intel Core i5 | 0 | 0 | In |
| 2 | HP | Notebook | 8 | Intel HD Graphics 620 | No OS | 1.86 | 30636.0000 | 0 | 0 | 141.211998 | Intel Core i5 | 0 | 256 | In |
| 3 | Apple | Ultrabook | 16 | AMD Radeon Pro 455 | macOS | 1.83 | 135195.3360 | 0 | 1 | 220.534624 | Intel Core i7 | 0 | 512 | AM |
| 4 | Apple | Ultrabook | 8 | Intel Iris Plus Graphics 650 | macOS | 1.37 | 96095.8080 | 0 | 1 | 226.983005 | Intel Core i5 | 0 | 256 | In |

Tr. [637].

```
df[ "Gpu brand" ].value_counts().reset_index()
```

Out [627]:

| Gpu brand | count |
|-----------|--------|
| 0 | Intel |
| 1 | Nvidia |
| 2 | AMD |
| 3 | ARM |

Tn [628]:

```
df[df["Gpu brand"] == 'ABM']
```

Out[628]:

| Company | TypeName | Ram | Gpu | OpSys | Weight | Price | Touchscreen | Ips | ppi | Cpu brand | HDD | SSD | Gi br. |
|---------|----------|-----|-----|-------|--------|-------|-------------|-----|-----|-----------|-----|-----|--------|
|---------|----------|-----|-----|-------|--------|-------|-------------|-----|-----|-----------|-----|-----|--------|

| Company | Type Name | Ram | Gpu | OpSys | Weight | Price | Touchscreen | Ips | ppi | Cpu brand | HDD | SSD | Gf brad |
|---------|-----------|--------------------|-----|---------------|-----------|-------|-------------|-----|-----|--------------------|-----|-----|---------|
| 1191 | Samsung | 2 in 1 Convertible | 4 | Mali T860 MP4 | Chrome OS | 1.15 | 35111.52 | 1 | 1 | 234.5074 Processor | 0 | 0 | AF |

In [629]:

```
df = df[df["Gpu brand"] != 'ARM']
```

In [630]:


```
df["Gpu brand"].value_counts().reset_index()
```

Out[630]:

| Gpu brand count | | |
|-----------------|--------|-----|
| 0 | Intel | 722 |
| 1 | Nvidia | 400 |
| 2 | AMD | 180 |

In [631]:

```
sns.barplot(x="Gpu brand", y="Price", data=df)
plt.title("GPU brand and it's cost")
plt.xlabel("Brand")
plt.ylabel("Cost")
plt.show()
```




In [632]:

```
sns.barplot(x="Gpu brand", y="Price", data=df, estimator=np.median)
plt.title("GPU brand and it's cost")
plt.xlabel("Brand")
plt.ylabel("Cost")
plt.show()
```

GPU brand and it's cost





In [633] :

```
#Drop 'Gpu' column since it is not required now:  
df.drop(columns="Gpu", inplace=True)
```

In [634] :

```
df.head()
```

Out [634] :

| | Company | TypeName | Ram | OpSys | Weight | Price | Touchscreen | Ips | ppi | Cpu brand | HDD | SSD | Gpu brand | |
|---|---------|-----------|-----|-------|--------|-------------|-------------|-----|-----|------------|---------------|-----|-----------|-------|
| 0 | Apple | Ultrabook | 8 | macOS | 1.37 | 71378.6832 | | 0 | 1 | 226.983005 | Intel Core i5 | 0 | 128 | Intel |
| 1 | Apple | Ultrabook | 8 | macOS | 1.34 | 47895.5232 | | 0 | 0 | 127.677940 | Intel Core i5 | 0 | 0 | Intel |
| 2 | HP | Notebook | 8 | No OS | 1.86 | 30636.0000 | | 0 | 0 | 141.211998 | Intel Core i5 | 0 | 256 | Intel |
| 3 | Apple | Ultrabook | 16 | macOS | 1.83 | 135195.3360 | | 0 | 1 | 220.534624 | Intel Core i7 | 0 | 512 | AMD |
| 4 | Apple | Ultrabook | 8 | macOS | 1.37 | 96095.8080 | | 0 | 1 | 226.983005 | Intel Core i5 | 0 | 256 | Intel |

In [] :

In [] :

Now we will do EDA On 'OpSys'(i.e. Operating System) and 'Weight' columns:

In [635] :

```
df[ "OpSys" ].value_counts().reset_index()
```

Out [635] :

| OpSys | count |
|---------|-------|
| Windows | 1000 |
| macOS | 500 |
| Ubuntu | 100 |
| Linux | 50 |
| Others | 100 |

| | OpSys | count |
|---|--------------|-------|
| 0 | Windows 10 | 1072 |
| 1 | No OS | 66 |
| 2 | Linux | 62 |
| 3 | Windows 7 | 45 |
| 4 | Chrome OS | 26 |
| 5 | macOS | 13 |
| 6 | Mac OS X | 8 |
| 7 | Windows 10 S | 8 |
| 8 | Android | 2 |

In [636]:

```
df["Weight"].value_counts().reset_index()
```

Out[636]:

| | Weight | count |
|-----|--------|-------|
| 0 | 2.200 | 126 |
| 1 | 2.100 | 58 |
| 2 | 2.000 | 45 |
| 3 | 2.400 | 44 |
| 4 | 2.300 | 41 |
| ... | ... | ... |
| 166 | 0.990 | 1 |
| 167 | 2.591 | 1 |
| 168 | 2.210 | 1 |
| 169 | 2.191 | 1 |
| 170 | 2.340 | 1 |

171 rows × 2 columns

In [637]:

```
def cat_os(inp):
    if inp == "Windows 10" or inp == "Windows 7" or inp == "Windows 10 S":
        return "Windows"
    elif inp == "macOS" or inp == "Mac OS X":
        return "Mac"
    else:
        return "Others/No OS/Linux"

df["os"] = df["OpSys"].apply(cat_os)
```

In [638]:

```
df.head()
```

Out[638]:

| | Company | TypeName | Ram | OpSys | Weight | Price | Touchscreen | Ips | ppi | Cpu brand | HDD | SSD | Gpu brand |
|---|---------|-----------|-----|-------|--------|------------|-------------|-----|------------|---------------|-----|-----|-----------|
| 0 | Apple | Ultrabook | 8 | macOS | 1.37 | 71378.6832 | 0 | 1 | 226.983005 | Intel Core i5 | 0 | 128 | Intel |
| 1 | Apple | Ultrabook | 8 | macOS | 1.34 | 47895.5232 | 0 | 0 | 127.677940 | Intel Core i5 | 0 | 0 | Intel |

| 2 | Company | HP | TypeName | Notebook | Ram | 8 | OpSys | No OS | Weight | 1.86 | Price | 30636.0000 | Touchscreen | 0 | Ips | 0 | ppi | 141.211998 | Cpu brand | Intel Core i5 | HDD | 0 | SSD | 256 | Gpu brand | AMD | Others OS/Linux |
|---|---------|-----------|----------|----------|------|-------------|-------|-------|------------|---------------|-------|------------|-------------|---|-----|---|-----|------------|-----------|---------------|-----|---|-----|-----|-----------|-----|-----------------|
| 3 | Apple | Ultrabook | 16 | macOS | 1.83 | 135195.3360 | 0 | 1 | 220.534624 | Intel Core i7 | 0 | 512 | AMD | 1 | | | | | | | | | | | | | |
| 4 | Apple | Ultrabook | 8 | macOS | 1.37 | 96095.8080 | 0 | 1 | 226.983005 | Intel Core i5 | 0 | 256 | Intel | 1 | | | | | | | | | | | | | |

In [639]:

```
df.drop(columns=[ "OpSys" ], inplace=True)
```

In [640]:


```
df.head()
```

Out[640]:

| | Company | TypeName | Ram | Weight | Price | Touchscreen | Ips | ppi | Cpu brand | HDD | SSD | Gpu brand | os |
|---|---------|-----------|-----|--------|-------------|-------------|-----|------------|---------------|-----|-----|-----------|--------------------|
| 0 | Apple | Ultrabook | 8 | 1.37 | 71378.6832 | 0 | 1 | 226.983005 | Intel Core i5 | 0 | 128 | Intel | Mac |
| 1 | Apple | Ultrabook | 8 | 1.34 | 47895.5232 | 0 | 0 | 127.677940 | Intel Core i5 | 0 | 0 | Intel | Mac |
| 2 | HP | Notebook | 8 | 1.86 | 30636.0000 | 0 | 0 | 141.211998 | Intel Core i5 | 0 | 256 | Intel | Others/No OS/Linux |
| 3 | Apple | Ultrabook | 16 | 1.83 | 135195.3360 | 0 | 1 | 220.534624 | Intel Core i7 | 0 | 512 | AMD | Mac |
| 4 | Apple | Ultrabook | 8 | 1.37 | 96095.8080 | 0 | 1 | 226.983005 | Intel Core i5 | 0 | 256 | Intel | Mac |

In [641]:


```
sns.barplot(y="os", x="Price", data=df)
plt.show()
```



Now we will do EDA on 'Weight' column:


In [642]:

```
#Distribution of all values in the 'Weight' column:  
sns.distplot(df["Weight"])  
plt.show()
```



In [643]:

```
sns.scatterplot(x="Weight", y="Price", data=df)  
plt.show()
```



In [644]:

```
#Correlation of all numeric columns with respect to 'Price' column:  
df.corr(numeric_only=True) ["Price"].reset_index()
```

Out[644]:

Out [644]:

| index | Price |
|-------|----------------------|
| 0 | Ram 0.742905 |
| 1 | Weight 0.209867 |
| 2 | Price 1.000000 |
| 3 | Touchscreen 0.192917 |
| 4 | Lps 0.253320 |
| 5 | ppi 0.475368 |
| 6 | HDD -0.096891 |
| 7 | SSD 0.670660 |

In []:

```
[ ]:
```

In []:

```
[ ]:
```

In [645]:

```
#Correlation of all columns with respect to all other columns:  
df.corr(numeric_only=True)
```

Out[645]:

| | Ram | Weight | Price | Touchscreen | Lps | ppi | HDD | SSD |
|-------------|----------|-----------|-----------|-------------|-----------|-----------|-----------|-----------|
| Ram | 1.000000 | 0.383362 | 0.742905 | 0.118875 | 0.207949 | 0.305688 | 0.095808 | 0.603379 |
| Weight | 0.383362 | 1.000000 | 0.209867 | -0.293004 | 0.018643 | -0.321883 | 0.514147 | -0.063818 |
| Price | 0.742905 | 0.209867 | 1.000000 | 0.192917 | 0.253320 | 0.475368 | -0.096891 | 0.670660 |
| Touchscreen | 0.118875 | -0.293004 | 0.192917 | 1.000000 | 0.148026 | 0.458571 | -0.208766 | 0.257577 |
| Lps | 0.207949 | 0.018643 | 0.253320 | 0.148026 | 1.000000 | 0.299142 | -0.093588 | 0.225311 |
| ppi | 0.305688 | -0.321883 | 0.475368 | 0.458571 | 0.299142 | 1.000000 | -0.294698 | 0.509437 |
| HDD | 0.095808 | 0.514147 | -0.096891 | -0.208766 | -0.093588 | -0.294698 | 1.000000 | -0.400750 |
| SSD | 0.603379 | -0.063818 | 0.670660 | 0.257577 | 0.225311 | 0.509437 | -0.400750 | 1.000000 |

Above result on a heatmap:


In [646]:

```
#Correlation of all columns with respect to all other columns:  
sns.heatmap(df.corr(numeric_only=True))
```

Out[646]:

<Axes: >





In []:

In []:

In []:

Now we will do EDA on 'Price' column:

In [647]:

```
X = df.drop(columns=["Price"])
y = np.log(df["Price"])
```

In [648]:

X

Out[648]:

| | Company | TypeName | Ram | Weight | Touchscreen | Ips | ppi | Cpu brand | HDD | SSD | Gpu brand | os |
|------|---------|--------------------|-----|--------|-------------|-----|------------|-----------------------|------|-----|-----------|--------------------|
| 0 | Apple | Ultrabook | 8 | 1.37 | 0 | 1 | 226.983005 | Intel Core i5 | 0 | 128 | Intel | Mac |
| 1 | Apple | Ultrabook | 8 | 1.34 | 0 | 0 | 127.677940 | Intel Core i5 | 0 | 0 | Intel | Mac |
| 2 | HP | Notebook | 8 | 1.86 | 0 | 0 | 141.211998 | Intel Core i5 | 0 | 256 | Intel | Others/No OS/Linux |
| 3 | Apple | Ultrabook | 16 | 1.83 | 0 | 1 | 220.534624 | Intel Core i7 | 0 | 512 | AMD | Mac |
| 4 | Apple | Ultrabook | 8 | 1.37 | 0 | 1 | 226.983005 | Intel Core i5 | 0 | 256 | Intel | Mac |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 1298 | Lenovo | 2 in 1 Convertible | 4 | 1.80 | 1 | 1 | 157.350512 | Intel Core i7 | 0 | 128 | Intel | Windows |
| 1299 | Lenovo | 2 in 1 Convertible | 16 | 1.30 | 1 | 1 | 276.053530 | Intel Core i7 | 0 | 512 | Intel | Windows |
| 1300 | Lenovo | Notebook | 2 | 1.50 | 0 | 0 | 111.935204 | Other Intel Processor | 0 | 0 | Intel | Windows |
| 1301 | HP | Notebook | 6 | 2.19 | 0 | 0 | 100.454670 | Intel Core i7 | 1000 | 0 | AMD | Windows |
| 1302 | Asus | Notebook | 4 | 2.20 | 0 | 0 | 100.454670 | Other Intel Processor | 500 | 0 | Intel | Windows |

In [649]:

y

Out[649]:

```
0      11.175755
1      10.776777
2      10.329931
3      11.814476
4      11.473101
...
1298   10.433899
1299   11.288115
1300   9.409283
1301   10.614129
1302   9.886358
Name: Price, Length: 1302, dtype: float64
```

In [650]:

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.15,random_state=2)

from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline

#load all models to check which algorithm will perform best (since we dont know which it would be):
from sklearn.linear_model import LinearRegression, Ridge, Lasso
from sklearn.neighbors import KNeighborsRegressor
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor, AdaBoostRegressor, ExtraTreesRegressor
from sklearn.svm import SVR
from xgboost import XGBRegressor
from sklearn.metrics import r2_score,mean_absolute_error

#Create a column transformer:
from sklearn.preprocessing import OneHotEncoder
step1 = ColumnTransformer(
    transformers=[
        ("col_tnf", OneHotEncoder(sparse_output=False, drop="first"), [0, 1, 7, 10, 11])
    ], remainder='passthrough')

step2 = LinearRegression()

pipe = Pipeline([
    ('step1', step1),
    ('step2', step2)
])

pipe.fit(X_train,y_train)

y_pred = pipe.predict(X_test)

print('R2 score = ',r2_score(y_test,y_pred))
print('MAE     = ',mean_absolute_error(y_test,y_pred))
```

R2 score = 0.8073277450155012

MAE = 0.21017827953019166

Now start testing all algorithms and see which algorithm gives the best result:

LINEAR REGRESSION:

In [651]:

```
# Same preprocessing: OneHotEncode specific columns, keep the rest
step1 = ColumnTransformer(
    transformers=[
        ('col_tnf', OneHotEncoder(sparse_output=False, drop='first'), [0, 1, 7, 10, 11])
    ],
    remainder='passthrough'
)

# Step 2: Ordinary Linear Regression (no regularization)
step2 = LinearRegression()

# Create the pipeline
pipe = Pipeline([
    ('step1', step1),
    ('step2', step2)
])

# Train the model
pipe.fit(X_train, y_train)

# Predictions
y_pred = pipe.predict(X_test)

# Evaluation
print('R2 score:', r2_score(y_test, y_pred))
print('MAE:', mean_absolute_error(y_test, y_pred))
```

R2 score: 0.8073277450155012
MAE: 0.21017827953019166

RIDGE REGRESSION:

In [652]:

```
step1 = ColumnTransformer(
    transformers=[
        ('col_tnf', OneHotEncoder(sparse_output=False, drop='first'), [0, 1, 7, 10, 11])
    ],
    remainder='passthrough' # keep all other columns unchanged
)

# Step 2: Ridge regression with alpha=10
step2 = Ridge(alpha=10)

# Create the pipeline
pipe = Pipeline([
    ('step1', step1),
    ('step2', step2)
])

# Split data (assuming you have X_train, X_test, y_train, y_test already)
# If not, use train_test_split
pipe.fit(X_train, y_train)

# Predictions
y_pred = pipe.predict(X_test)

# Evaluation
print('R2 score:', r2_score(y_test, y_pred))
print('MAE:', mean_absolute_error(y_test, y_pred))
```

R2 score: 0.8127331033739459
MAE: 0.20926802210371442

LASSO REGRESSION:

In [653]:

```
step1 = ColumnTransformer(
    transformers=[
        ('col_tnf', OneHotEncoder(sparse_output=False, drop='first'), [0, 1, 7, 10, 11])
    ],
    remainder='passthrough' # numerical columns pass through unchanged
)

# Step 2: Lasso regression with alpha=0.001
step2 = Lasso(alpha=0.001)

# Create the pipeline
pipe = Pipeline([
    ('step1', step1),
    ('step2', step2)
])

# Train the model
pipe.fit(X_train, y_train)

# Predictions on test set
y_pred = pipe.predict(X_test)

# Evaluation
print('R2 score:', r2_score(y_test, y_pred))
print('MAE:', mean_absolute_error(y_test, y_pred))
```

R2 score: 0.8071853947620582
MAE: 0.21114361575113458

KNN:

In [654]:

```
step1 = ColumnTransformer(
    transformers=[
        ('col_tnf', OneHotEncoder(sparse_output=False, drop='first'), [0, 1, 7, 10, 11])
    ],
    remainder='passthrough' # numerical columns pass through as is
)

# Step 2: K-Nearest Neighbors Regressor with 3 neighbors
step2 = KNeighborsRegressor(n_neighbors=3)

# Create the pipeline
pipe = Pipeline([
    ('step1', step1),
    ('step2', step2)
])

# Train the model
pipe.fit(X_train, y_train)

# Predictions on test set
y_pred = pipe.predict(X_test)

# Evaluation (optional but recommended)
print('R2 score:', r2_score(y_test, y_pred))
print('MAE:', mean_absolute_error(y_test, y_pred))
```

R2 score: 0.8031008164264897
MAE: 0.19268746498695286

DECISION TREE:

In [655]:

```
step1 = ColumnTransformer(
```

```

step1 = ColumnTransformer(
    transformers=[
        ('col_tnf', OneHotEncoder(sparse_output=False, drop='first'), [0, 1, 7, 10, 11])
    ],
    remainder='passthrough' # numerical columns pass through as is
)

step2 = DecisionTreeRegressor(max_depth=8)

pipe = Pipeline([
    ('step1', step1),
    ('step2', step2)
])

pipe.fit(X_train, y_train)
y_pred = pipe.predict(X_test)

print('R2 score:', r2_score(y_test, y_pred))
print('MAE:', mean_absolute_error(y_test, y_pred))

```

R2 score: 0.8426084411425243
MAE: 0.18107155164530095

SVM:

In [656]:

```

step1 = ColumnTransformer(
    transformers=[
        ('col_tnf', OneHotEncoder(sparse_output=False, drop='first'), [0, 1, 7, 10, 11])
    ],
    remainder='passthrough' # numerical features pass through as is
)

step2 = SVR(kernel='rbf', C=10000, epsilon=0.1)

# Create the pipeline
pipe = Pipeline([
    ('step1', step1),
    ('step2', step2)
])

# Train the model
pipe.fit(X_train, y_train)

# Predictions on test set
y_pred = pipe.predict(X_test)

# Evaluation
print('R2 score:', r2_score(y_test, y_pred))
print('MAE:', mean_absolute_error(y_test, y_pred))

```

R2 score: 0.8083168388457612
MAE: 0.20239400567814725

RANDOM FOREST:

In [657]:

```

step1 = ColumnTransformer(
    transformers=[
        ('col_tnf', OneHotEncoder(sparse_output=False, drop='first'), [0, 1, 7, 10, 11])
    ],
    remainder='passthrough' # numerical features pass through as is
)

# Step 2: Random Forest Regressor with specified hyperparameters
step2 = RandomForestRegressor(
    n_estimators=100,

```

```

        random_state=3,
        max_samples=0.5,
        max_features=0.75,
        max_depth=15
    )

# Create the pipeline
pipe = Pipeline([
    ('step1', step1),
    ('step2', step2)
])

# Train the model
pipe.fit(X_train, y_train)

# Predictions on test set
y_pred = pipe.predict(X_test)

# Evaluation
print('R2 score:', r2_score(y_test, y_pred))
print('MAE:', mean_absolute_error(y_test, y_pred))

```

R2 score: 0.8873402378382488
MAE: 0.15860130110457718

EXTRATREES:

In [658]:

```

step1 = ColumnTransformer(
    transformers=[
        ('col_tnf', OneHotEncoder(sparse_output=False, drop='first'), [0, 1, 7, 10, 11])
    ],
    remainder='passthrough' # numerical features pass through as is
)

# Step 2: ExtraTrees Regressor with specified hyperparameters
step2 = ExtraTreesRegressor(
    n_estimators=100,
    random_state=3,
    max_samples=0.5,
    max_features=0.75,
    max_depth=15,
    bootstrap=True # required when using max_samples < 1.0
)

# Create the pipeline
pipe = Pipeline([
    ('step1', step1),
    ('step2', step2)
])

# Train the model
pipe.fit(X_train, y_train)

# Predictions on test set
y_pred = pipe.predict(X_test)

# Evaluation
print('R2 score:', r2_score(y_test, y_pred))
print('MAE:', mean_absolute_error(y_test, y_pred))

```

R2 score: 0.8850720167552375
MAE: 0.16154538000217084

ADABOOST:

In [659]:

1 2 3 4 5 6 7 8 9

```

step1 = ColumnTransformer(
    transformers=[
        ('col_tnf', OneHotEncoder(sparse_output=False, drop='first'), [0, 1, 7, 10, 11])
    ],
    remainder='passthrough' # numerical features pass through as is
)

step2 = AdaBoostRegressor(n_estimators=15, learning_rate=1.0)

# Create the pipeline
pipe = Pipeline([
    ('step1', step1),
    ('step2', step2)
])

# Train the model
pipe.fit(X_train, y_train)

# Predictions on test set
y_pred = pipe.predict(X_test)

# Evaluation
print('R2 score:', r2_score(y_test, y_pred))
print('MAE:', mean_absolute_error(y_test, y_pred))

```

R2 score: 0.8020965600051384
MAE: 0.2211495197113982

GRADIENT BOOST:

In [660]:

```

step1 = ColumnTransformer(transformers=[
    ('col_tnf', OneHotEncoder(sparse_output=False, drop='first'), [0,1,7,10,11])
], remainder='passthrough')

step2 = GradientBoostingRegressor(n_estimators=500)

pipe = Pipeline([
    ('step1', step1),
    ('step2', step2)
])

pipe.fit(X_train, y_train)

y_pred = pipe.predict(X_test)

print('R2 score', r2_score(y_test, y_pred))
print('MAE', mean_absolute_error(y_test, y_pred))

```

R2 score 0.8815819094071498
MAE 0.15966219871732804

XG BOOST:

In [661]:

```

step1 = ColumnTransformer(transformers=[
    ('col_tnf', OneHotEncoder(sparse_output=False, drop='first'), [0,1,7,10,11])
], remainder='passthrough')

step2 = XGBRegressor(n_estimators=45, max_depth=5, learning_rate=0.5)

pipe = Pipeline([
    ('step1', step1),
    ('step2', step2)
])

```

```

pipe.fit(X_train, y_train)

y_pred = pipe.predict(X_test)

print('R2 score', r2_score(y_test, y_pred))
print('MAE', mean_absolute_error(y_test, y_pred))

```

R2 score 0.8771404806375557
MAE 0.16262936288951352

VOTING REGRESSOR:

In [662]:

```

step1 = ColumnTransformer(transformers=[
    ('col_tnf', OneHotEncoder(sparse_output=False, drop='first'), [0,1,7,10,11])
], remainder='passthrough')

# Fixed: Added bootstrap=True for both RF and ET when using max_samples
rf = RandomForestRegressor(
    n_estimators=350,
    random_state=3,
    max_samples=0.5,           # subsample 50% of data per tree
    max_features=0.75,
    max_depth=15,
    bootstrap=True             # REQUIRED when max_samples is set
)

gbdt = GradientBoostingRegressor(n_estimators=100, max_features=0.5)

xgb = XGBRegressor(n_estimators=25, learning_rate=0.3, max_depth=5)

et = ExtraTreesRegressor(
    n_estimators=100,
    random_state=3,
    max_samples=0.5,           # subsample 50% of data per tree
    max_features=0.75,
    max_depth=15,
    bootstrap=True             # REQUIRED when max_samples is set
)

step2 = VotingRegressor(
    estimators=[('rf', rf), ('gbdt', gbdt), ('xgb', xgb), ('et', et)],
    weights=[5, 1, 1, 1]
)

pipe = Pipeline([
    ('step1', step1),
    ('step2', step2)
])

pipe.fit(X_train, y_train)
y_pred = pipe.predict(X_test)

print('R2 score:', r2_score(y_test, y_pred))
print('MAE:', mean_absolute_error(y_test, y_pred))

```

R2 score: 0.8906299333066894
MAE: 0.15721405706647654

STACKING:

In [663]:

```

from sklearn.ensemble import VotingRegressor, StackingRegressor

step1 = ColumnTransformer(transformers=[
    ('col_tnf', OneHotEncoder(sparse_output=False, drop='first'), [0,1,7,10,11])
])

```

```

], remainder='passthrough')

estimators = [
    ('rf', RandomForestRegressor(n_estimators=350, random_state=3, max_samples=0.5, max_
features=0.75, max_depth=15, bootstrap=True)),
    ('gbdt', GradientBoostingRegressor(n_estimators=100, max_features=0.5)),
    ('xgb', XGBRegressor(n_estimators=25, learning_rate=0.3, max_depth=5))
]

step2 = StackingRegressor(
    estimators=estimators,
    final_estimator=Ridge(alpha=100)
)

pipe = Pipeline([
    ('step1', step1),
    ('step2', step2)
])

pipe.fit(X_train, y_train)

y_pred = pipe.predict(X_test)

# Optional evaluation
from sklearn.metrics import r2_score, mean_absolute_error
print('R2 score:', r2_score(y_test, y_pred))
print('MAE:', mean_absolute_error(y_test, y_pred))

```

R2 score: 0.8788531484462075
MAE: 0.166733576119004

In []:

In []:

CONCLUSION FROM THE ML ALGORITHMS]

RandomForest performs the best hence we will use this model in our website which we are going to deploy for predicting Laptop Prices based on Laptop's configurations.

We will also export our dataframe which will be used in the dropdowns of our website"

In [664]:

```
df.head()
```

Out[664]:

| | Company | TypeName | Ram | Weight | Price | Touchscreen | Ips | ppi | Cpu brand | HDD | SSD | Gpu brand | os | |
|---|---------|-----------|-----|--------|-------------|-------------|-----|-----|------------|---------------|-----|-----------|-------|--------------------|
| 0 | Apple | Ultrabook | 8 | 1.37 | 71378.6832 | | 0 | 1 | 226.983005 | Intel Core i5 | 0 | 128 | Intel | Mac |
| 1 | Apple | Ultrabook | 8 | 1.34 | 47895.5232 | | 0 | 0 | 127.677940 | Intel Core i5 | 0 | 0 | Intel | Mac |
| 2 | HP | Notebook | 8 | 1.86 | 30636.0000 | | 0 | 0 | 141.211998 | Intel Core i5 | 0 | 256 | Intel | Others/No OS/Linux |
| 3 | Apple | Ultrabook | 16 | 1.83 | 135195.3360 | | 0 | 1 | 220.534624 | Intel Core i7 | 0 | 512 | AMD | Mac |
| 4 | Apple | Ultrabook | 8 | 1.37 | 96095.8080 | | 0 | 1 | 226.983005 | Intel Core i5 | 0 | 256 | Intel | Mac |

In [665]:

```
import pickle
pickle.dump(df, open('df.pkl', 'wb'))
pickle.dump(pipe, open('pipe.pkl', 'wb'))
```

In []: