

Laptops Segmentation into separate clusters using KMeans and Testing claims using Hypothesis Testing

In [33]:

```
import pandas as pd
import numpy as np

import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.preprocessing import (
    LabelEncoder,
    StandardScaler,
    OneHotEncoder
)

from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline

#Ignore all incoming warnings (if any):
from warnings import filterwarnings as f
f("ignore")

df = pd.read_csv("laptopPrice.csv")
```

In [34]:

```
print(df.head())
```

	brand	processor_brand	processor_name	processor_gnrtn	ram_gb	ram_type	\
0	ASUS	Intel	Core i3	10th	4 GB	DDR4	
1	Lenovo	Intel	Core i3	10th	4 GB	DDR4	
2	Lenovo	Intel	Core i3	10th	4 GB	DDR4	
3	ASUS	Intel	Core i5	10th	8 GB	DDR4	
4	ASUS	Intel	Celeron Dual	Not Available	4 GB	DDR4	

	ssd	hdd	os	os_bit	graphic_card_gb	weight	warranty	\
0	0 GB	1024 GB	Windows	64-bit	0 GB	Casual	No warranty	
1	0 GB	1024 GB	Windows	64-bit	0 GB	Casual	No warranty	
2	0 GB	1024 GB	Windows	64-bit	0 GB	Casual	No warranty	
3	512 GB	0 GB	Windows	32-bit	2 GB	Casual	No warranty	
4	0 GB	512 GB	Windows	64-bit	0 GB	Casual	No warranty	

	Touchscreen	msoffice	Price	rating	Number of Ratings	Number of Reviews
0	No	No	34649	2 stars	3	0
1	No	No	38999	3 stars	65	5
2	No	No	39999	3 stars	8	1
3	No	No	69990	3 stars	0	0
4	No	No	26990	3 stars	0	0

In [35]:

```
print(df.shape)
```

(823, 19)

In [36]:

```
print(df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 823 entries, 0 to 822
Data columns (total 19 columns):
#   Column                                Non-Null Count  Dtype
0   brand                                823 non-null    object
1   processor_brand                      823 non-null    object
2   processor_name                      823 non-null    object
3   processor_gnrtn                     823 non-null    object
4   ram_gb                              823 non-null    object
5   ram_type                            823 non-null    object
6   ssd                                 823 non-null    object
7   hdd                                 823 non-null    object
8   os                                   823 non-null    object
9   os_bit                              823 non-null    object
10  graphic_card_gb                     823 non-null    object
11  weight                              823 non-null    object
12  warranty                             823 non-null    object
13  Touchscreen                         823 non-null    object
14  msoffice                            823 non-null    object
15  Price                               823 non-null    object
16  rating                              823 non-null    object
17  Number of Ratings                   823 non-null    object
18  Number of Reviews                   823 non-null    object
```

	column	non-null count	dtype
0	brand	823 non-null	object
1	processor_brand	823 non-null	object
2	processor_name	823 non-null	object
3	processor_gnrtn	823 non-null	object
4	ram_gb	823 non-null	object
5	ram_type	823 non-null	object
6	ssd	823 non-null	object
7	hdd	823 non-null	object
8	os	823 non-null	object
9	os_bit	823 non-null	object
10	graphic_card_gb	823 non-null	object
11	weight	823 non-null	object
12	warranty	823 non-null	object
13	Touchscreen	823 non-null	object
14	msoffice	823 non-null	object
15	Price	823 non-null	int64
16	rating	823 non-null	object
17	Number of Ratings	823 non-null	int64
18	Number of Reviews	823 non-null	int64

dtypes: int64(3), object(16)

memory usage: 122.3+ KB

None

In [37]:

```
df.isnull().sum()
```

Out[37]:

```
brand          0
processor_brand 0
processor_name  0
processor_gnrtn 0
ram_gb         0
ram_type       0
ssd            0
hdd            0
os             0
os_bit         0
graphic_card_gb 0
weight         0
warranty       0
Touchscreen    0
msoffice       0
Price          0
rating         0
Number of Ratings 0
Number of Reviews 0
dtype: int64
```

In [38]:

```
# percentage missing

(df.isnull().sum()/len(df))*100
```

Out[38]:

```
brand          0.0
processor_brand 0.0
processor_name  0.0
processor_gnrtn 0.0
ram_gb         0.0
ram_type       0.0
ssd            0.0
hdd            0.0
os             0.0
os_bit         0.0
graphic_card_gb 0.0
weight         0.0
warranty       0.0
dtype: float64
```

```
touchscreen      0.0
msoffice         0.0
Price            0.0
rating           0.0
Number of Ratings 0.0
Number of Reviews 0.0
dtype: float64
```

In [39]:

```
df.duplicated().sum()
```

Out[39]:

```
np.int64(21)
```

In [40]:

```
# remove duplicates
df.drop_duplicates(inplace=True)
```

In [41]:

```
df.columns = df.columns.str.lower()

df.head()
```

```
# RAM conversion
```

```
df["ram_gb"] = (
    df["ram_gb"]
    .str.replace(" GB", "")
    .astype(int)
)
```

```
# SSD conversion
```

```
df["ssd"] = (
    df["ssd"]
    .str.replace(" GB", "")
    .astype(int)
)
```

```
# HDD conversion
```

```
df["hdd"] = (
    df["hdd"]
    .str.replace(" GB", "")
    .astype(int)
)
```

```
# Graphic card
```

```
df["graphic_card_gb"] = (
    df["graphic_card_gb"]
    .str.extract("(\d+) ")
    .fillna(0)
    .astype(int)
)
```

```
# Rating conversion
```

```

df["rating"] = (
    df["rating"]
    .str.replace(" stars", "", regex=False)
    .str.replace(" star", "", regex=False)
    .astype(int)
)

# OS bit

df["os_bit"] = (
    df["os_bit"]
    .str.replace("-bit", "")
    .astype(int)
)

df.head()

# =====
# 7. CREATE NEW FEATURES
# =====

# Total storage

df["total_storage"] = (
    df["ssd"] + df["hdd"]
)

# Storage type

df["has_ssd"] = np.where(
    df["ssd"] > 0, 1, 0
)

# GPU availability

df["has_gpu"] = np.where(
    df["graphic_card_gb"] > 0, 1, 0
)

# Laptop performance score

df["performance_score"] = (
    df["ram_gb"] * 2 +
    df["processor_name"].map({
        "Core i3": 3,
        "Core i5": 5,
        "Core i7": 7,
        "Core i9": 9,
        "Ryzen 3": 3,
        "Ryzen 5": 5,
        "Ryzen 7": 7,
        "M1": 8,
        "Celeron Dual": 2,
        "Pentium Quad": 2
    }).fillna(0)
    +
    df["graphic_card_gb"]
)

```

```

# customer popularity

df["popularity_score"] = (
    df["number of ratings"] +
    df["number of reviews"]
)

# Total Storage:
df["total_storage"] = (
    df["ssd"] + df["hdd"]
)

#GPU availability:
df["has_gpu"] = np.where(
    df["graphic_card_gb"]>0,
    1,
    0
)

df.drop(
    [
        "processor_name"
    ],
    axis=1,
    inplace=True
)

```

Testing claims using Hypothesis Testing

Problem statement: Test whether the average price of laptops is different from ₹60,000. Large sample size justifies z-test. (one-sample ztest for means)

In [42]:

```

from statsmodels.stats.weightstats import ztest

z_stat, p_value = ztest(df["price"], value=60000)

if p_value < 0.05:
    print("Reject H0: Average price significantly different from ₹60,000.")
else:
    print("Fail to reject H0.")

```

Reject H0: Average price significantly different from ₹60,000.

Problem statement: Test whether average price of HP differs from Lenovo (large samples) (two-sample ztest).

In [43]:

```

price_hp = df[df["brand"] == "HP"]["price"]
price_lenovo = df[df["brand"] == "Lenovo"]["price"]

z_stat, p_value = ztest(price_hp, price_lenovo)

if p_value < 0.05:
    print("Reject H0: Significant price difference between HP and Lenovo.")
else:
    print("Fail to reject H0.")

```

Fail to reject H0.

Problem statement: Is the proportion of laptops with SSD greater than 70%? (One-sample z-test for proportion)

In [44]:

```
from statsmodels.stats.proportion import proportions_ztest

count = df["has_ssd"].sum()

z_stat, p_value = proportions_ztest(count, len(df), value=0.70, alternative="larger")

if p_value < 0.05:
    print("Reject H0: SSD proportion > 70%.")
else:
    print("Fail to reject H0.")
```

Reject H0: SSD proportion > 70%.

Problem statement: Is the proportion of laptops with SSD less than 40%? (One-sample z-test for proportion)

In [45]:

```
count = df["has_ssd"].sum()
nobs = len(df)
value = 0.40

z_stat, p_value = proportions_ztest(count, nobs, value=value, alternative="smaller")

if p_value < 0.05:
    print("Reject H0: SSD proportion > 70%.")
else:
    print("Fail to reject H0.")
```

Fail to reject H0.

Problem statement: Average rating of Apple laptops differs from 3.5 stars? (One-sample t-test)

In [46]:

```
from scipy.stats import ttest_1samp

t_stat, p_value = ttest_1samp(df[df["brand"]=="APPLE"]["rating"], popmean=3.5)

if p_value < 0.05:
    print("Reject H0: Average rating ≠ 3.5.")
else:
    print("Fail to reject H0.")
```

Reject H0: Average rating ≠ 3.5.

Problem statement: Is there any price difference between laptops with GPU vs without GPU? (Independent two-sample t-test)

In [47]:

```
from scipy.stats import ttest_ind

price_with_gpu = df[df["has_gpu"] == 1]["price"]
price_without = df[df["has_gpu"] == 0]["price"]

t_stat, p_value = ttest_ind(price_with_gpu, price_without, equal_var=False)

if p_value < 0.05:
    print("Reject H0: Significant price difference.")
else:
    print("Fail to reject H0.")
```

Reject H0: Significant price difference.

Problem statement: Mean difference between number of ratings and number of reviews ≠ 0? (Paired t-test)

In [48]:

```
from scipy.stats import ttest_rel

t_stat, p_value = ttest_rel(df["number of ratings"], df["number of reviews"])

if p_value < 0.05:
    print("Reject H0: Mean difference ≠ 0.")
else:
    print("Fail to reject H0.")
```

Reject H0: Mean difference ≠ 0.

Problem statement: Does average price differ by processor brand (Intel, AMD, others)? (One-way ANOVA)

In [49]:

```
from scipy.stats import f_oneway

groups = [group["price"].values for name, group in df.groupby("processor_brand") if len(
group) > 1]

f_stat, p_value = f_oneway(*groups)

print(f"F: {f_stat:.4f}, p-value: {p_value:.4f}")
if p_value < 0.05:
    print("Reject H0: At least one processor brand has different avg price.")
else:
    print("Fail to reject H0.")
```

F: 30.5862, p-value: 0.0000

Reject H0: At least one processor brand has different avg price.

Problem statement: Effects of brand (ASUS, HP, Lenovo) and GPU presence on price, including interaction. (Two-way ANOVA)

In [50]:

```
import statsmodels.api as sm
from statsmodels.formula.api import ols

df_2w = df[df["brand"].isin(["ASUS", "HP", "Lenovo"])].copy()
model = ols('price ~ C(brand) * C(has_gpu)', data=df_2w).fit()
anova_table = sm.stats.anova_lm(model, typ=2)
print(anova_table)

print("\n\n\n")
print("The ANOVA results indicate that GPU availability is the most influential factor among the tested variables in determining laptop prices, while brand alone does not have a statistically significant impact.")
```

	sum_sq	df	F	PR(>F)
C(brand)	2.596207e+09	2.0	0.737480	4.788203e-01
C(has_gpu)	1.488046e+11	1.0	84.539020	9.198666e-19
C(brand):C(has_gpu)	1.403273e+10	2.0	3.986144	1.914482e-02
Residual	9.100176e+11	517.0	NaN	NaN

The ANOVA results indicate that GPU availability is the most influential factor among the tested variables in determining laptop prices, while brand alone does not have a statistically significant impact.

In []:

In []:

Segmented Products into different clusters using KMeans

Problem statement : "A laptop company wants to segment its laptops into different customer segments based on specifications, pricing, and customer feedback. Build a clustering model that groups laptops into meaningful categories such as Budget laptops, Professional laptops, Gaming laptops, and Premium laptops."

In [51]:

```
df.shape
```

Out[51]:

```
(802, 23)
```

In [52]:

```
cluster_df = df[
    [
        "price",
        "ram_gb",
        "graphic_card_gb",
        "total_storage",
        "performance_score",
        "rating",
        "number of ratings",
        "number of reviews",
        "brand",
        "processor_brand",
        "os"
    ]
]

cat_cols = cluster_df.select_dtypes(include="object").columns
num_cols = cluster_df.select_dtypes(exclude="object").columns

preprocessor = ColumnTransformer(
    transformers=[
        (
            "categorical",
            OneHotEncoder(handle_unknown="ignore"),
            cat_cols
        ),
        (
            "numeric",
            StandardScaler(),
            num_cols
        )
    ]
)

X = preprocessor.fit_transform(cluster_df)
```

In [53]:

```
from sklearn.cluster import KMeans

wcscs = []

for k in range(1,11):

    kmeans = KMeans(
        n_clusters=k,
        random_state=42,
        n_init=10
    )
```



```

kmeans.fit(X)

wcss.append(kmeans.inertia_)


plt.figure(figsize=(8,5))

plt.plot(
    range(1,11),
    wcss,
    marker="o"
)

plt.xlabel("Number of Clusters (K)")
plt.ylabel("WCSS")
plt.title("Elbow Method")

plt.show()

```



In [54]:

```

from sklearn.metrics import silhouette_score

for k in range(2,11):

    kmeans = KMeans(
        n_clusters=k,
        random_state=42,
        n_init=10
    )

    labels = kmeans.fit_predict(X)

    score = silhouette_score(X, labels)

    print(f"K = {k} --> Silhouette Score = {score:.4f}")

```

```

K = 2 --> Silhouette Score = 0.3039
K = 3 --> Silhouette Score = 0.3262
K = 4 --> Silhouette Score = 0.1927
K = 5 --> Silhouette Score = 0.2006
K = 6 --> Silhouette Score = 0.2071
K = 7 --> Silhouette Score = 0.2006
K = 8 --> Silhouette Score = 0.2006
K = 9 --> Silhouette Score = 0.2006
K = 10 --> Silhouette Score = 0.2006

```

```
K = / --> Silhouette Score = 0.2041
K = 8 --> Silhouette Score = 0.2084
K = 9 --> Silhouette Score = 0.2163
K = 10 --> Silhouette Score = 0.2091
```

In [55]:

```
kmeans = KMeans(
    n_clusters=3,
    random_state=42,
    n_init=10
)

clusters = kmeans.fit_predict(X)

cluster_df["cluster"] = clusters
```

In [56]:

```
pd.crosstab(
    cluster_df["cluster"],
    cluster_df["brand"]
)
```

Out[56]:

	brand	APPLE	ASUS	Avita	DELL	HP	Lenovo	MSI	acer
cluster									
0	9	175	5	134	95	117	27	40	
1	12	58	1	15	35	26	23	10	
2	2	10	0	0	6	1	0	1	

In [57]:

```
from sklearn.decomposition import PCA

pca = PCA(n_components=2)

X_pca = pca.fit_transform(X)

plt.figure(figsize=(10,6))


cluster_names = {
    0: "Premium / High-Performance Laptops",
    1: "Professional Laptops",
    2: "Budget & Mass-Market Laptops"
}

cluster_labels = pd.Series(clusters).map(cluster_names)

sns.scatterplot(
    x=X_pca[:,0],
    y=X_pca[:,1],
    hue=cluster_labels,
    palette="Set1"
)

plt.xlabel("PC1")
plt.ylabel("PC2")
```

```
plt.legend(title="Laptop Segment")
plt.title("Laptop Segmentation using KMeans and PCA")
plt.show()
```



Although the business initially expected four segments (Budget, Professional, Gaming, Premium), Silhouette Analysis indicated that the optimal number of clusters is 3 (score = 0.3278). Therefore, K-Means with K=3 was selected to obtain the most meaningful and well-separated customer segments.

From the above visualization we can say that Premium Laptops are the most dominate ones and the Budget-friendly laptops are very few. rofessional Laptops are present in a moderate quantity..

Now we will explore what type of information do each cluster holds and draw some insights based on it:

In [58]:

```
cluster_profile = cluster_df.groupby("cluster").agg({
    "price": "mean",
    "ram_gb": "mean",
    "graphic_card_gb": "mean",
    "total_storage": "mean",
    "performance_score": "mean",
    "rating": "mean",
    "number of ratings": "mean",
    "number of reviews": "mean"
}).round(2)

cluster_profile
```

Out[58]:

	price	ram_gb	graphic_card_gb	total_storage	performance_score	rating	number of ratings	number of reviews
cluster								
0	63649.39	6.57	0.86	651.91	18.28	3.58	198.92	24.28
1	123263.53	15.96	2.43	795.73	40.77	3.51	81.94	11.42
2	47465.95	6.40	0.90	678.40	17.65	3.80	5298.70	613.50

In [59]:

```
pd.crosstab(
    cluster_df["cluster"],
    cluster_df["brand"]
)
```

Out[59]:

brand	APPLE	ASUS	Avita	DELL	HP	Lenovo	MSI	acer
cluster								
0	9	175	5	134	95	117	27	40
1	12	58	1	15	35	26	23	10
2	2	10	0	0	6	1	0	1

In [60]:

```
pd.crosstab(
    cluster_df["cluster"],
    cluster_df["os"]
)
```

Out[60]:

os	DOS	Mac	Windows
cluster			
0	10	9	583
1	5	12	163
2	1	2	17

In []:

In []:

CONCLUSION: The K-Means algorithm clearly suggests that premium laptops are the most dominant segment, while budget-friendly laptops are comparatively fewer. Since the dataset contained multiple features, we analyzed the clusters using K-Means along with dimensionality reduction techniques such as PCA to visualize and understand the underlying patterns.

In []: